



SPAIN AI
[MÁLAGA]

Presencial



open
south
code

Panelistas:

Catedrático de
Inteligencia Artificial
de la UMA

EZEQUIEL LÓPEZ

Málaga AI en el Opensouthcode 2023

Panel de discusión:

LA REVOLUCIÓN OPEN SOURCE EN IA: STABLEDIFFUSION, LLAMA, SAM, DOLLY Y MÁS...



Moderado por
Adrián Tineo
City Leader Málaga AI

Viernes 9 Junio 18h.
La Térmica, Málaga
(Sala Benalmádena)

RESERVA EN OPENSOUTHCODE.ORG

Profesor de la UMA
e Investigador en IA
en Biomedicina

FRANCISCO J. VEREDAS

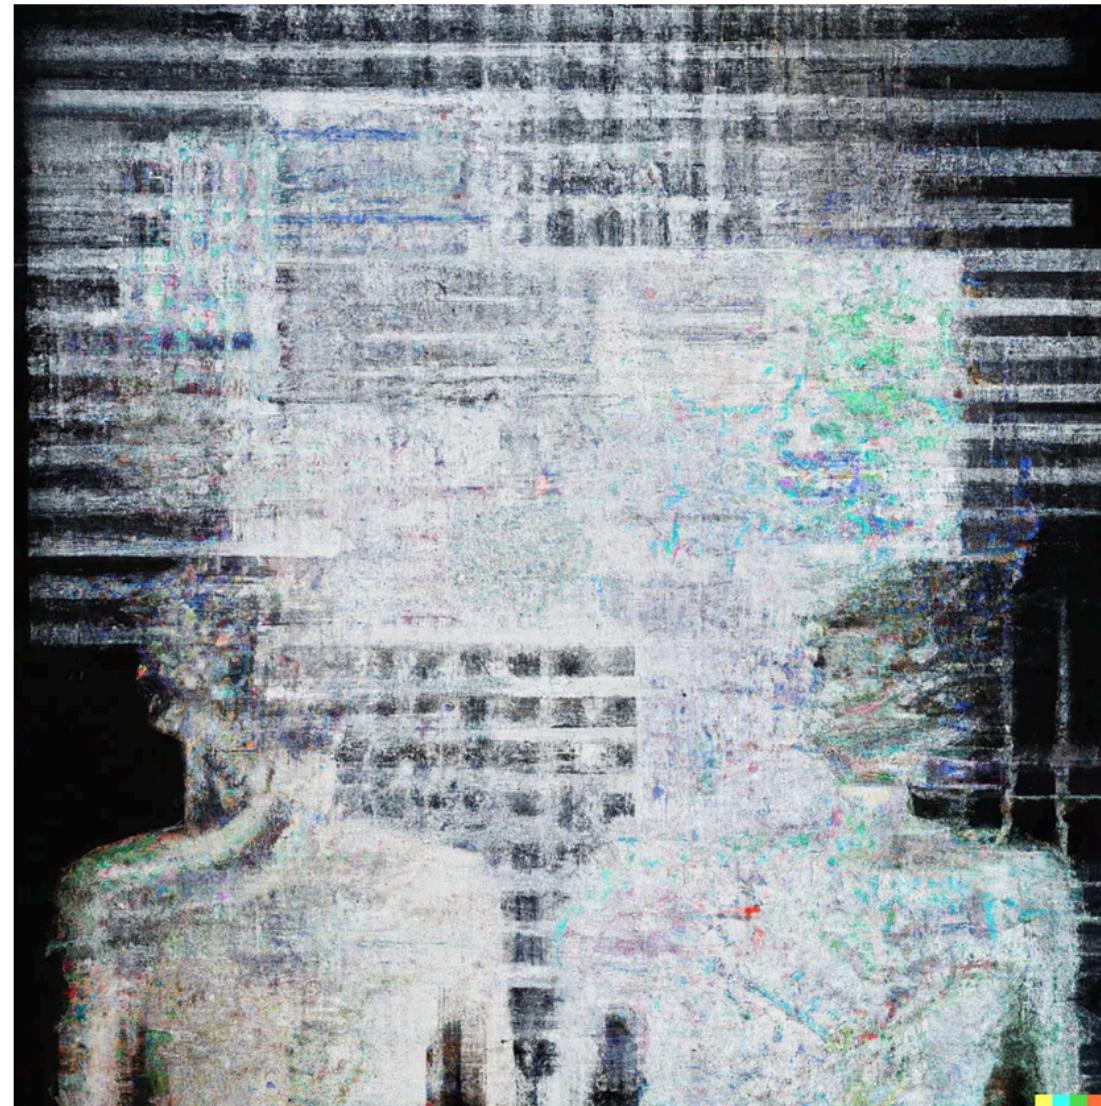
Síguenos en @AIMálaga



Dec 2022

A Smarter Robot

A new chatbot shows rapid advances in artificial intelligence.



An A.I. image generator produced this when given the prompt “A distributed linguistic superbrain that takes the form of an A.I. chatbot.” Kevin Roose, via DALL-E



By [German Lopez](#)

Dec. 8, 2022

“ChatGPT was a moment when a technology people had heard about finally became real to them”

<https://www.nytimes.com/2022/12/08/briefing/ai-chatgpt-openai.html>

Image generation

Anyone got access? 😢

Jul 2022

OpenAI Research ▾ Product ▾ Developers ▾ Safety Company ▾ Search Log in ▾ Sign up ▾

DALL-E now available in beta

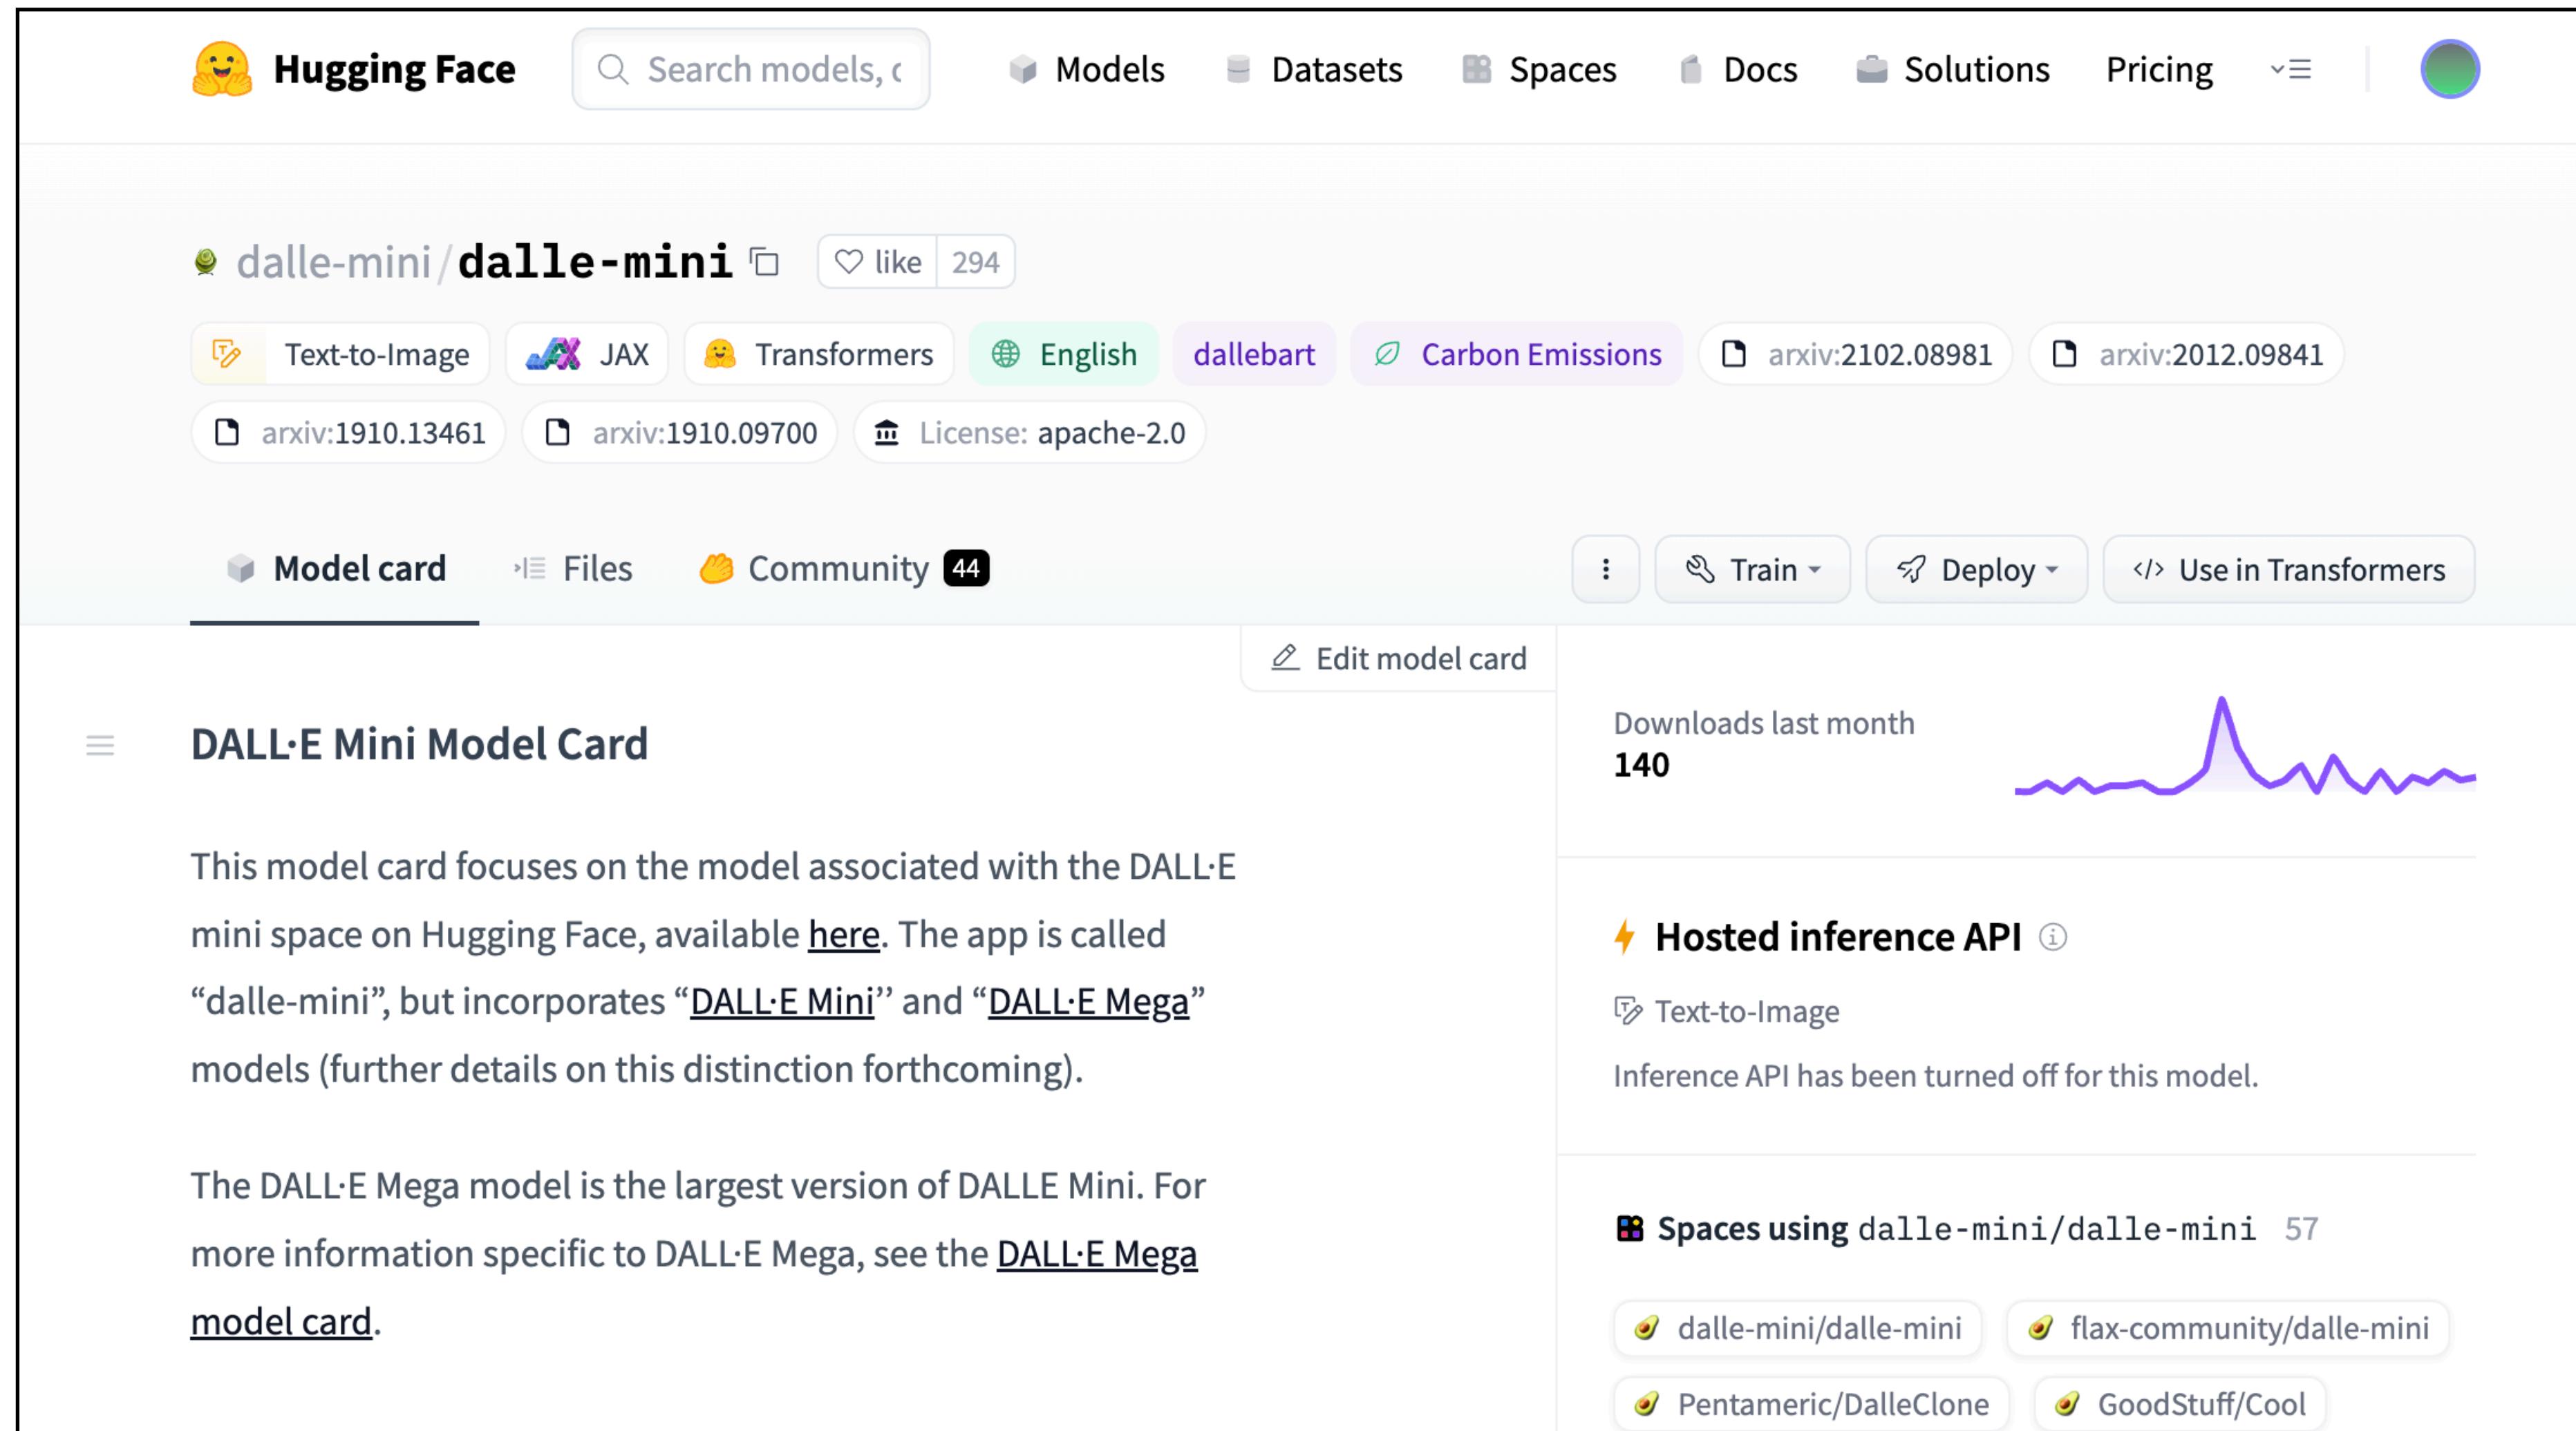
We'll invite 1 million people from our waitlist over the coming weeks. Users can create with DALL-E using free credits that refill every month, and buy additional credits in 115-generation increments for \$15.



Illustration: Ruby Chen & Justin Jay Wang x DALL-E

Here is access 😭

Jun 2022



The screenshot shows the Hugging Face platform interface for the model space `dalle-mini/dalle-mini`. The top navigation bar includes links for Models, Datasets, Spaces, Docs, Solutions, Pricing, and a user profile icon. A search bar is also present.

The main content area displays the `dalle-mini/dalle-mini` model card. Key details shown include:

- Tags:** Text-to-Image, JAX, Transformers, English, dallebart, Carbon Emissions.
- ArXiv IDs:** arxiv:2102.08981, arxiv:2012.09841, arxiv:1910.13461, arxiv:1910.09700.
- Licenses:** Apache-2.0.
- Downloads last month:** 140 (represented by a purple line chart).
- Hosted inference API:** Text-to-Image (Inference API has been turned off for this model).
- Spaces using dalle-mini/dalle-mini:** 57 (with examples: `dalle-mini/dalle-mini`, `flax-community/dalle-mini`, `Pentameric/DalleClone`, `GoodStuff/Cool`).

The `DALL·E Mini Model Card` section contains the following text:

This model card focuses on the model associated with the DALL·E mini space on Hugging Face, available [here](#). The app is called “dalle-mini”, but incorporates “[DALL·E Mini](#)” and “[DALL·E Mega](#)” models (further details on this distinction forthcoming).

The DALL·E Mega model is the largest version of DALLE Mini. For more information specific to DALL·E Mega, see the [DALL·E Mega model card](#).

<https://huggingface.co/dalle-mini/dalle-mini/tree/main>

Diffusion models go viral 😜

Jun 2022

Menú

El Confidencial

Iniciar sesión Suscríbete

Rajoy comiendo noodles: el programador detrás de la idea que ha vuelto loco a internet

El funcionamiento de Dall-E Mini es tan sencillo como teclear unas palabras y esperar a que genere imágenes relacionadas. Ni sus creadores se explican cómo, pero está por todas partes

DALL·E mini
AI model generating images from any prompt!

king juan carlos I painted as gta character

Run

Juan Carlos I en su versión del videojuego GTA, según Dall-E Mini.

Aug 2022

The model was trained on our 4,000 A100 Ezra-1 AI ultracluster over the last month as the first of a series of models exploring this and other approaches.

We have been testing the model at scale with over 10,000 beta testers that are creating 1.7 million images a day.



User generated image from Stable Diffusion Beta

Stable Diffusion changed everything... forever 😱





What does it take to learn how to draw?

Aug 2022

4.2 gigabytes.

That's the size of the model that has made this recent explosion possible.

4.2 gigabytes of floating points that somehow encode so much of what we know.

Yes, I'm waxing poetic here. No, I am not heralding the arrival of AGI, or our AI overlords. I am simply admiring the beauty of it, while it is fresh and new.

Because it won't be fresh and new for long. This thing I'm feeling is not much different from how I felt using email for the first time - "Grandma got my message already? In *Florida*? In seconds?" It was the nearest thing to magic my child-self had ever seen. Now email is the most boring and mundane part of my day.

The limits of “fair use”

Feb 2023

ARTIFICIAL INTELLIGENCE / TECH / LAW

Getty Images sues AI art generator Stable Diffusion in the US for copyright infringement



/ Getty Images has filed a case against Stability AI, alleging that the company copied 12 million images to train its AI model ‘without permission ... or compensation.’

By [James Vincent](#), a senior reporter who has covered AI, robotics, and more for eight years at The Verge.

Feb 6, 2023, 5:56 PM GMT+1 | □ [16 Comments](#) / [16 New](#)



An illustration from Getty Images' lawsuit, showing an original photograph and a similar image (complete with Getty Images watermark) generated by Stable Diffusion. Image: Getty Images

Text generation

Jun 2017

The Transformer architecture

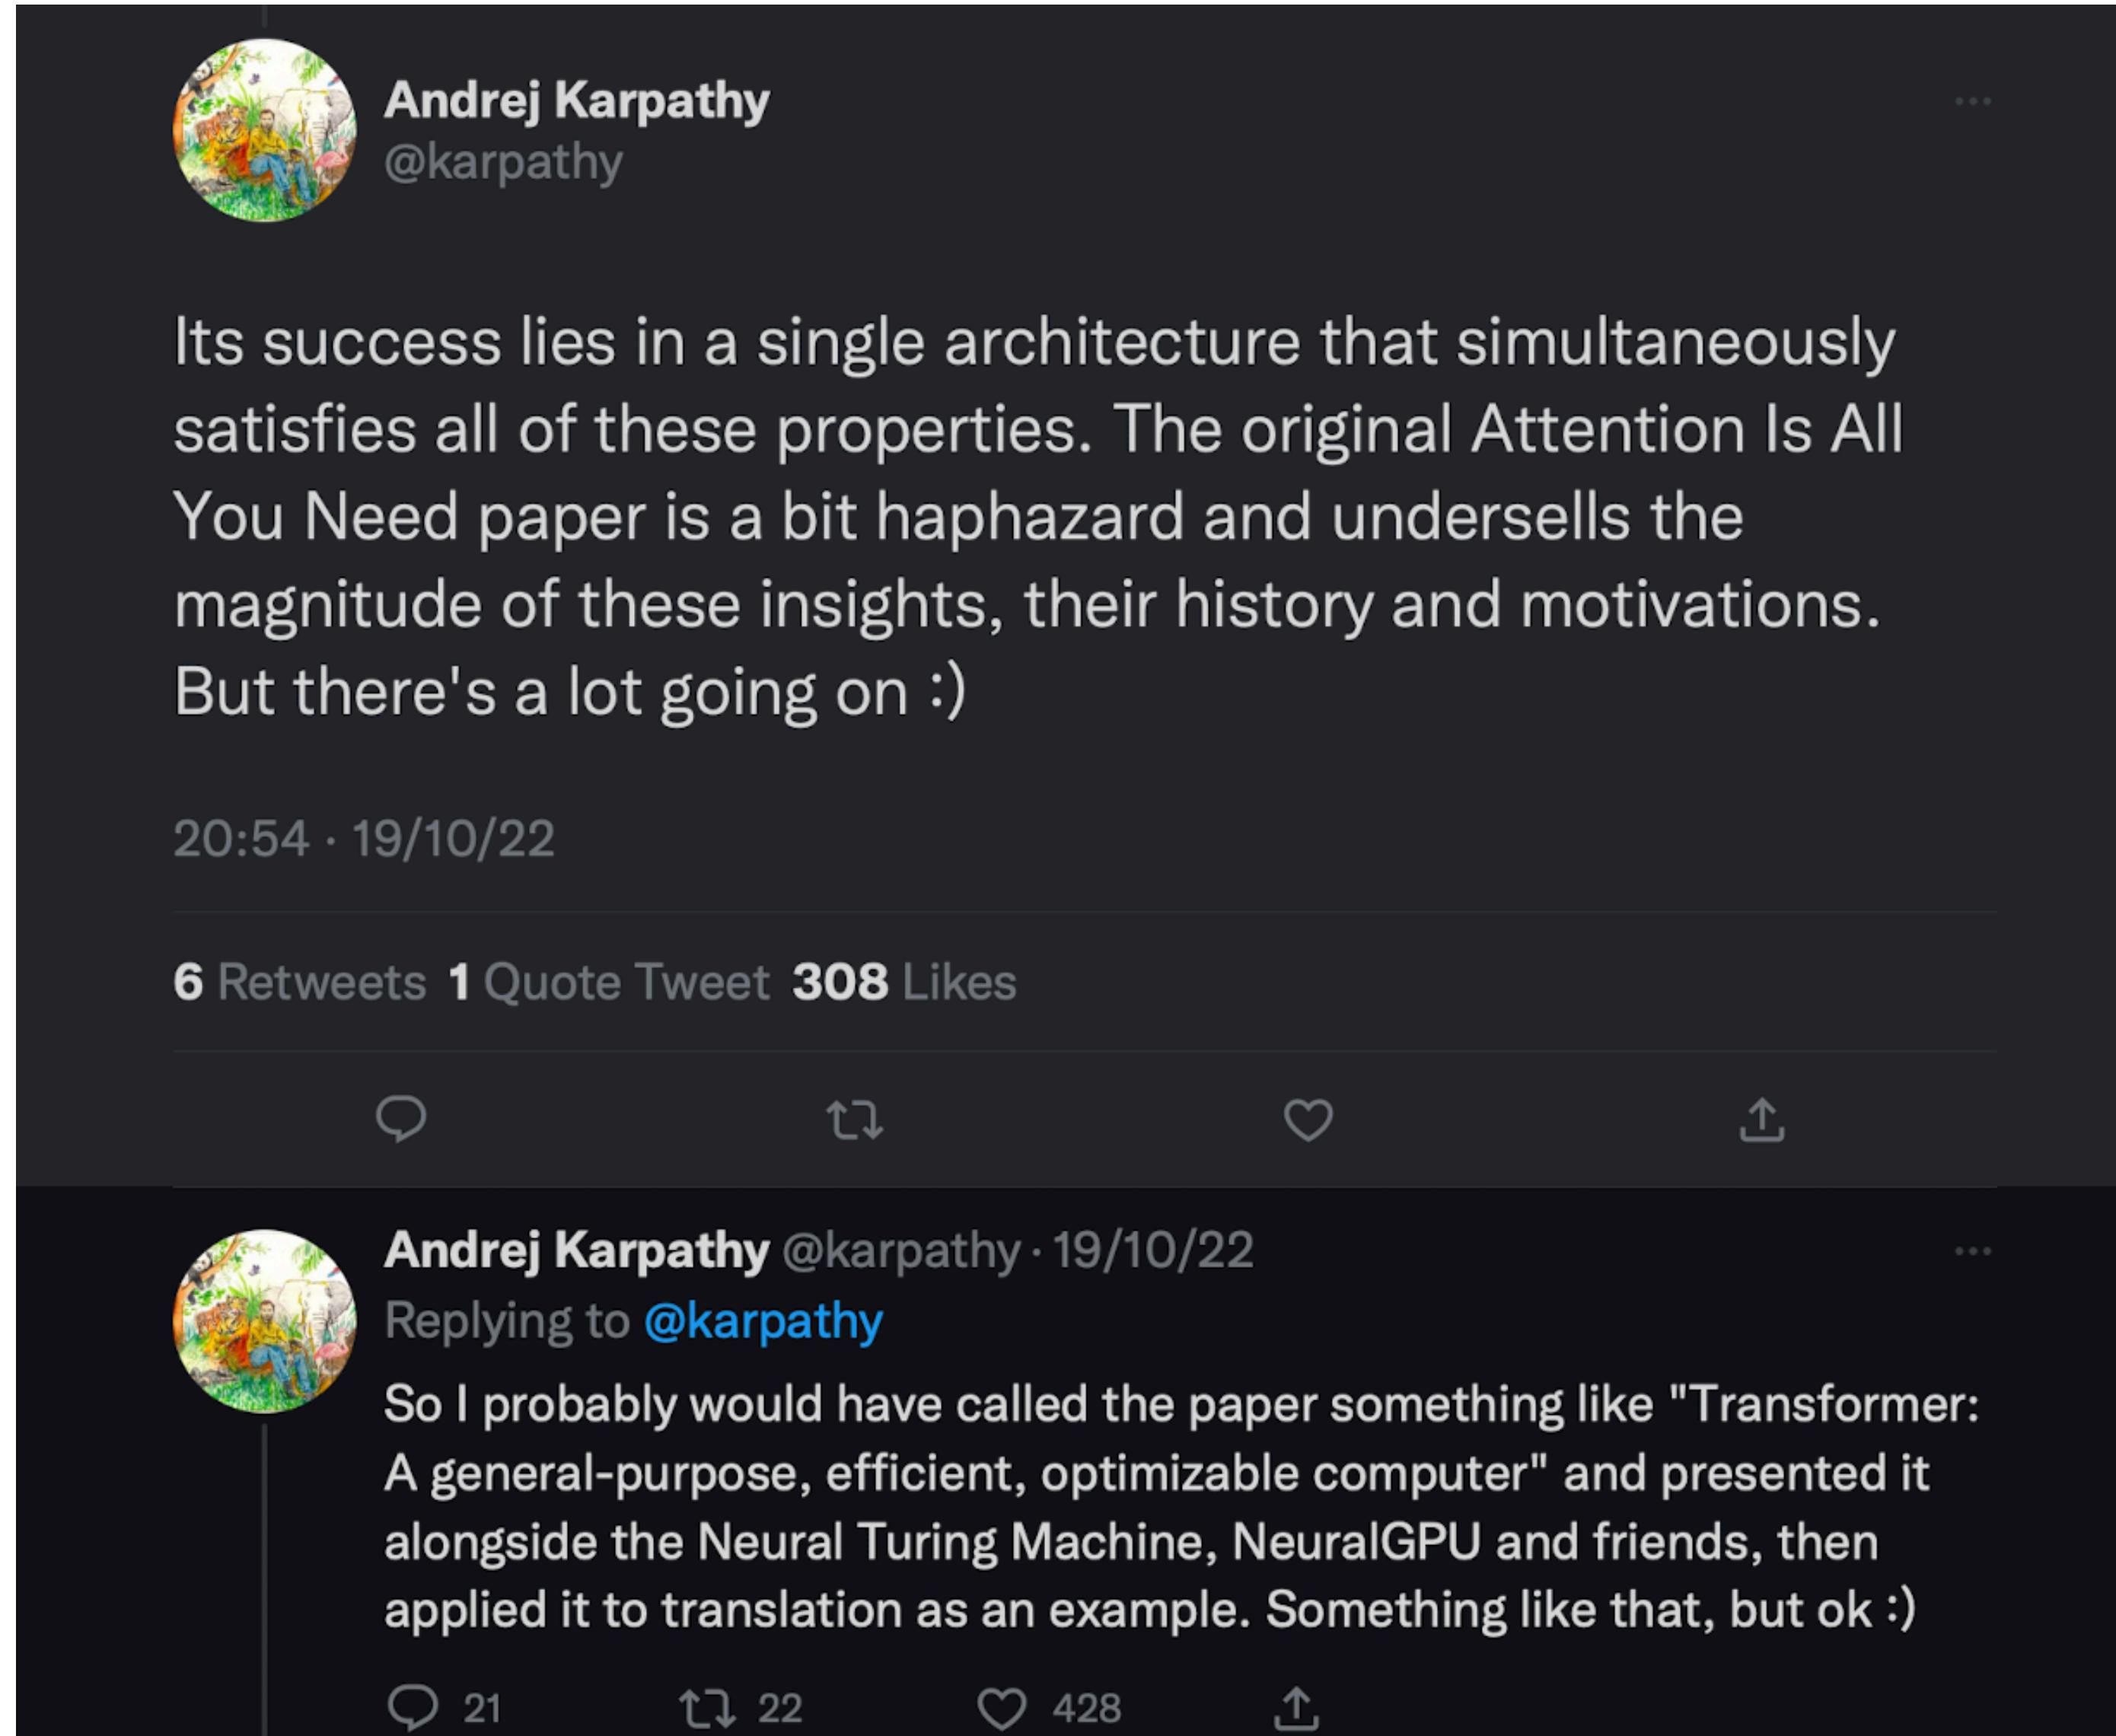
[Submitted on 12 Jun 2017 ([v1](#)), last revised 6 Dec 2017 (this version, v5)]

Attention Is All You Need

[Ashish Vaswani](#), [Noam Shazeer](#), [Niki Parmar](#), [Jakob Uszkoreit](#), [Llion Jones](#), [Aidan N. Gomez](#), [Łukasz Kaiser](#), [Illia Polosukhin](#)

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks in an encoder-decoder configuration. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

The Transformer architecture surpassed even its authors' expectations 😮



Its success lies in a single architecture that simultaneously satisfies all of these properties. The original Attention Is All You Need paper is a bit haphazard and undersells the magnitude of these insights, their history and motivations. But there's a lot going on :)

20:54 · 19/10/22

6 Retweets 1 Quote Tweet 308 Likes

Andrej Karpathy @karpathy · 19/10/22
Replying to @karpathy

So I probably would have called the paper something like "Transformer: A general-purpose, efficient, optimizable computer" and presented it alongside the Neural Turing Machine, NeuralGPU and friends, then applied it to translation as an example. Something like that, but ok :)

<https://twitter.com/karpathy/status/1582807372841373696>

175B parameters and \$4.6M to train 😂

Mar 2021

 OpenAI Research ▾ Product ▾ Developers ▾ Safety Company ▾ Search Log in ↗ Sign up ↗

GPT-3 powers the next generation of apps

Over 300 applications are delivering GPT-3-powered search, conversation, text completion, and other advanced AI features through our API.

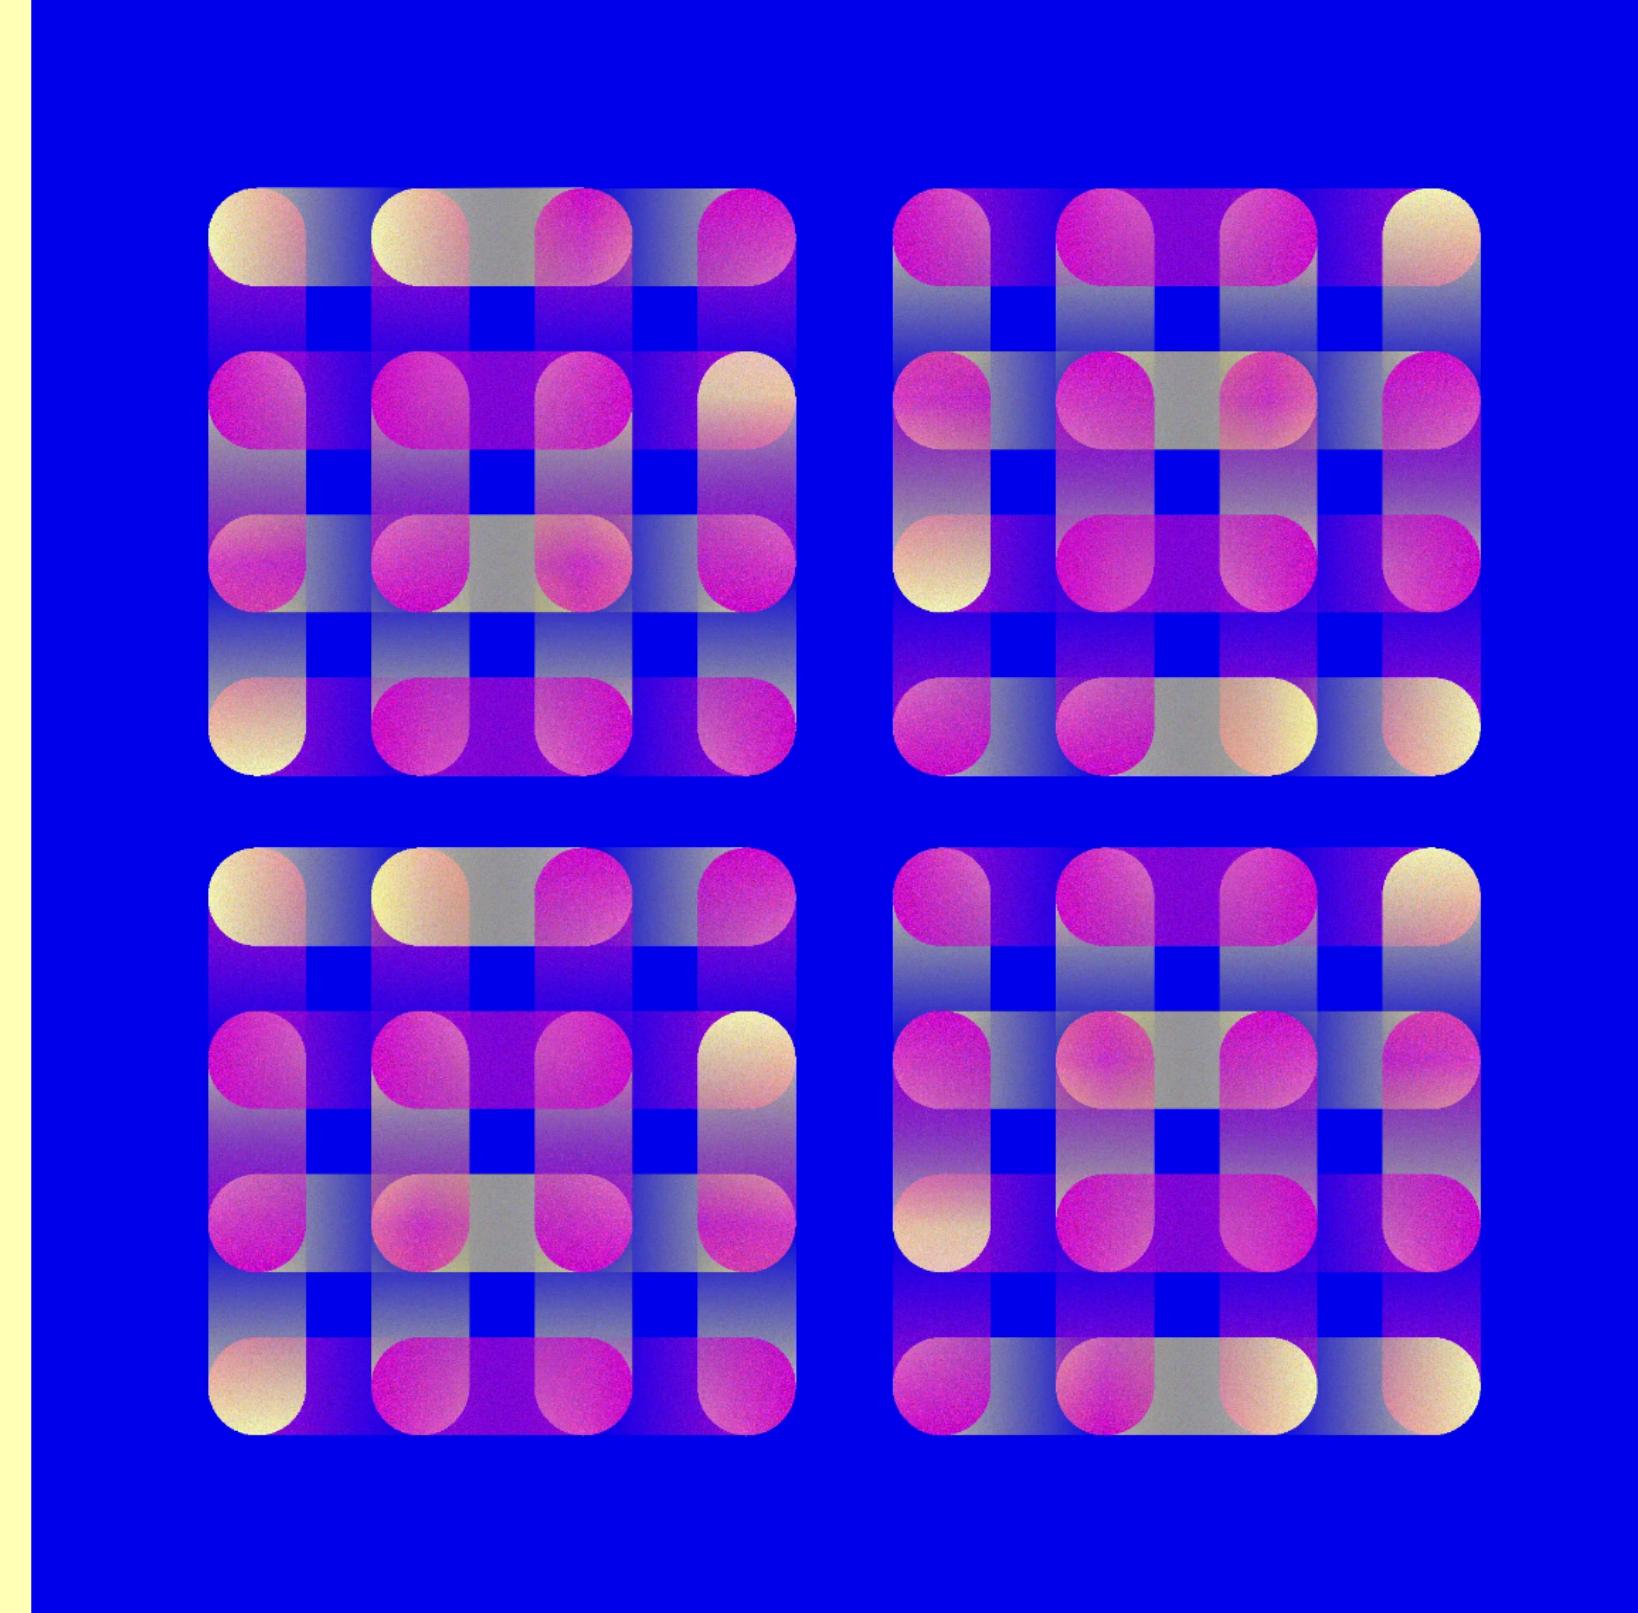
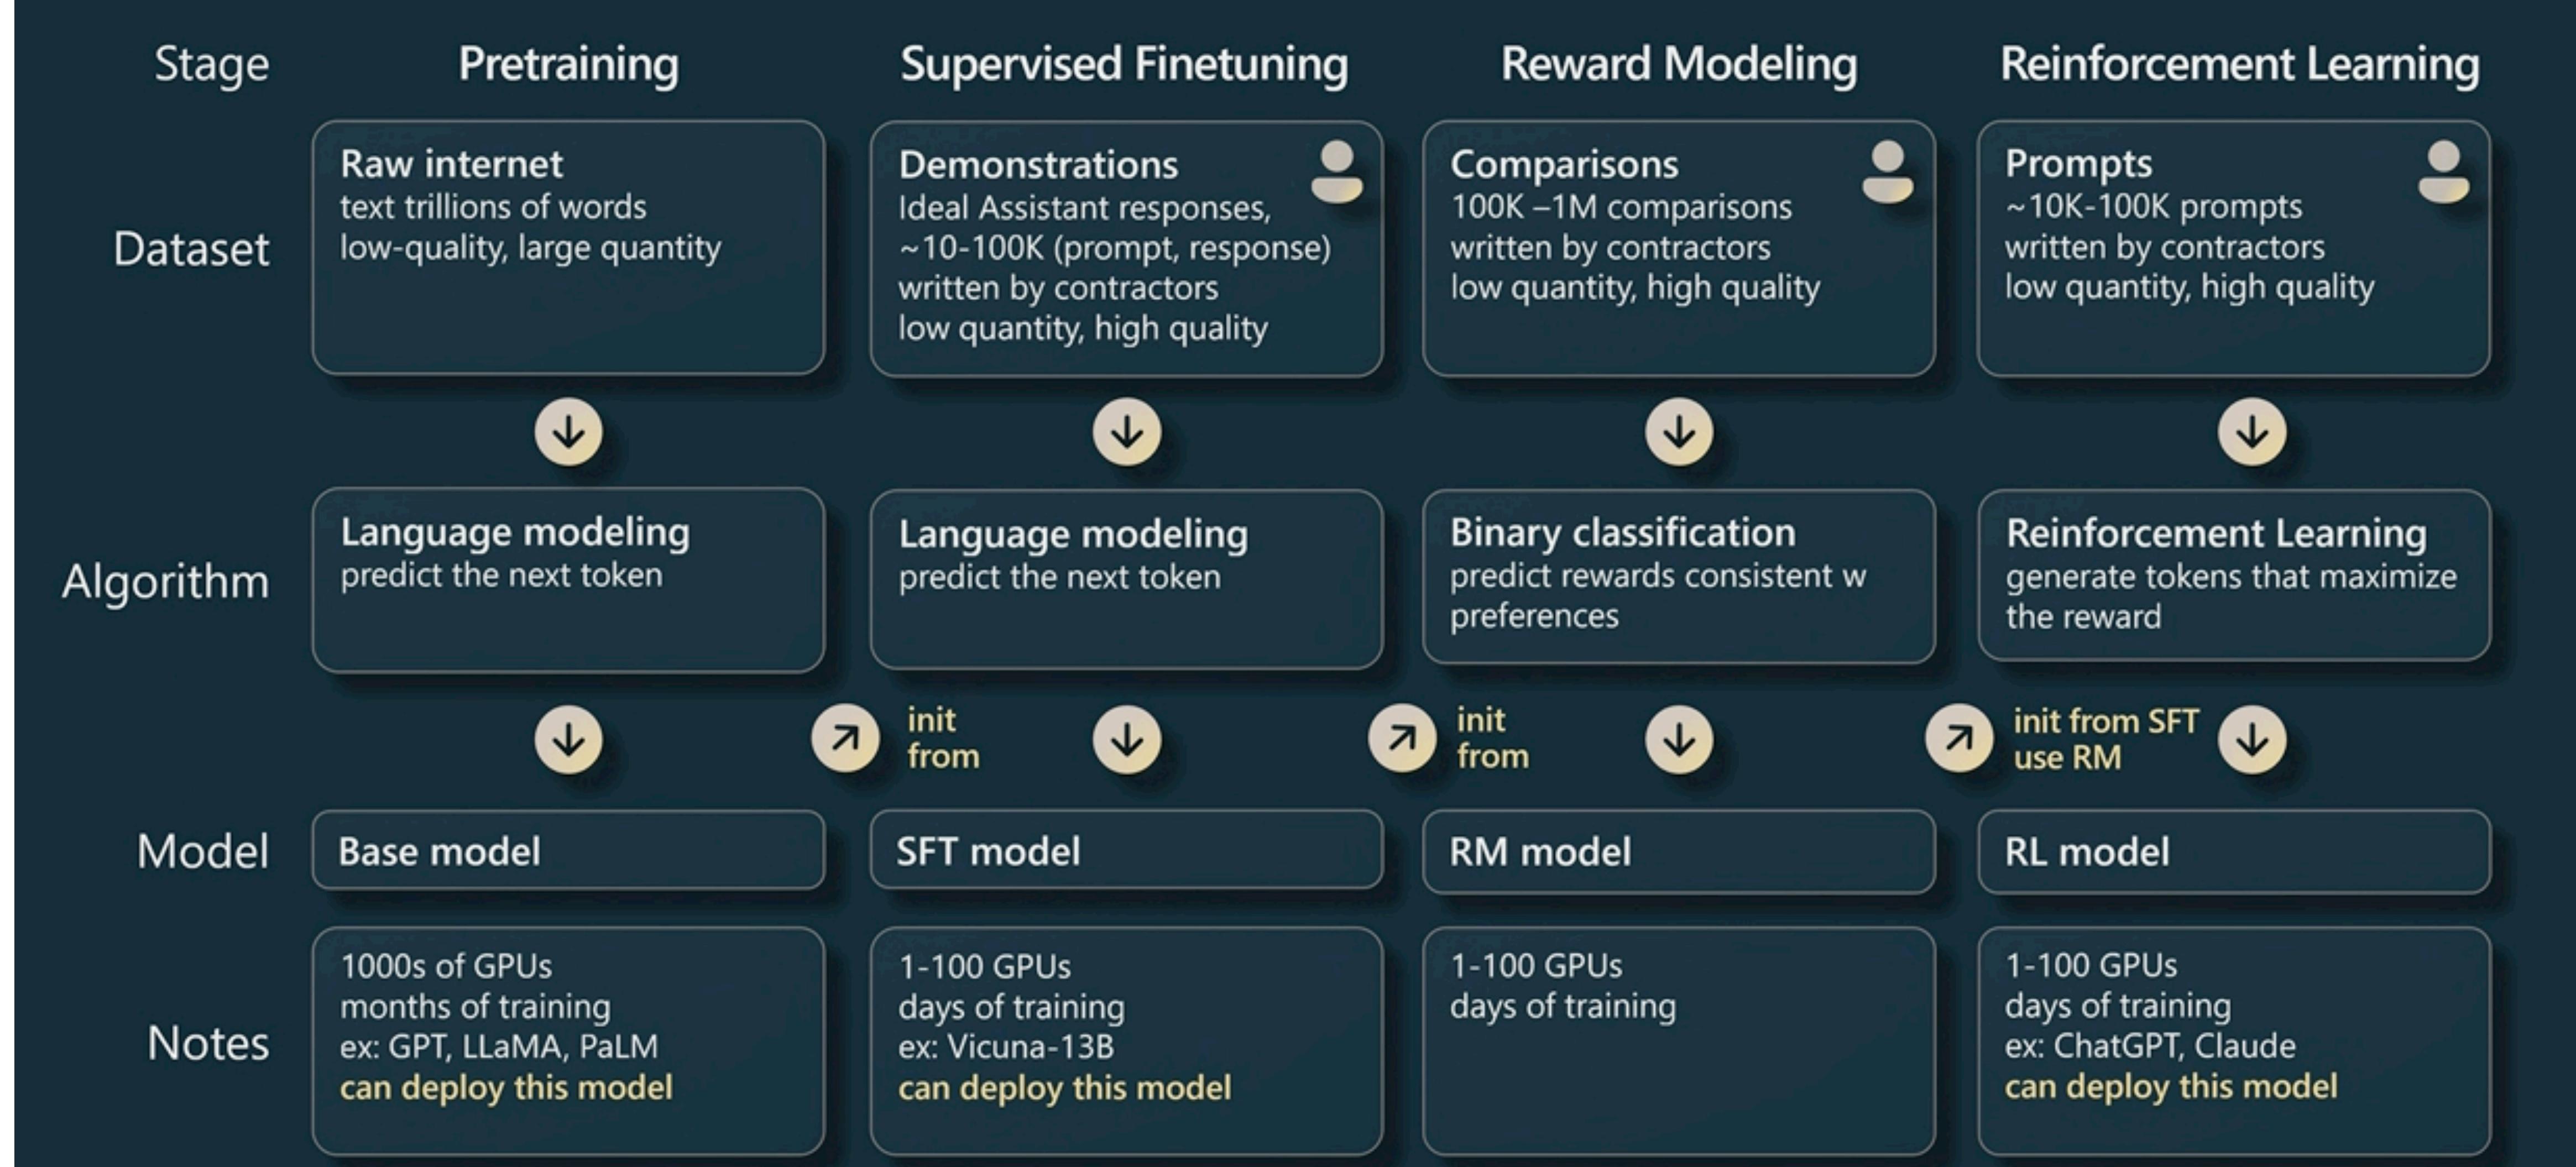


Illustration: Ruby Chen

GPT Assistant training pipeline





Open Source LLM Leaderboard

Jun 2023

Model	Revision	Average	ARC (25-shot)	HellaSwag (10-shot)
tiiuae/falcon-40b-instruct	main	63.2	61.6	84.4
tiiuae/falcon-40b	main	60.4	61.9	85.3
ausboss/llama-30b-supercot	main	59.8	58.5	82.9
llama-65b	main	58.3	57.8	84.2
MetaIX/GPT4-X-Alpasta-30b	main	57.9	56.7	81.4
Aeala/VicUnlocked-alpaca-30b	main	57.6	55	80.8
digitous/Alpacino30b	main	57.4	57.1	82.6
Aeala/GPT4-x-AlpacaDente2-30b	main	57.2	56.1	79.8
TheBloke/Wizard-Vicuna-13B-Uncensored-HF	main	57	53.6	79.6
TheBloke/dromedary-65b-lora-HF	main	57	57.8	80.8
llama-30b	main	56.9	57.1	82.6
elinas/llama-30b-hf-transformers-4.29	main	56.9	57.1	82.6
cyl/awsome-llama	main	56.8	54.4	79.7
openaccess-ai-collective/wizard-mega-13b	main	55.7	52.5	78.6

Big Tech & Open Source LLM Leaderboard

May 2023

Rank	Model	Elo Rating	Description
1	🟡 gpt-4	1225	ChatGPT-4 by OpenAI
2	🔴 claude-v1	1195	Claude by Anthropic
3	🟠 claude-instant-v1	1153	Claude Instant by Anthropic
4	gpt-3.5-turbo	1143	ChatGPT-3.5 by OpenAI
5	vicuna-13b	1054	a chat assistant fine-tuned from LLaMA on user-shared conversations by LMSYS
6	palm-2	1042	PaLM 2 for Chat (chat-bison@001) by Google
7	vicuna-7b	1007	a chat assistant fine-tuned from LLaMA on user-shared conversations by LMSYS
8	koala-13b	980	a dialogue model for academic research by BAIR
9	mpt-7b-chat	952	a chatbot fine-tuned from MPT-7B by MosaicML
10	fastchat-t5-3b	941	a chat assistant fine-tuned from FLAN-T5 by LMSYS
11	alpaca-13b	937	a model fine-tuned from LLaMA on instruction-following demonstrations by Stanford
12	RWKV-4-Raven-14B	928	an RNN with transformer-level LLM performance
13	oasst-pythia-12b	921	an Open Assistant for everyone by LAION
14	chatglm-6b	921	an open bilingual dialogue language model by Tsinghua University
15	stablelm-tuned-alpha-7b	882	Stability AI language models
16	dolly-v2-12b	866	an instruction-tuned open large language model by Databricks

Mar 2023

Ground-breaking foundation model



The Batch > AI & Society > Article

Runaway LLaMA

How Meta's LLaMA NLP model leaked

AI & Society NLP Transformer Bias Harm
Conversation Meta 4Chan Discord

Published Mar 15, 2023 Reading time 2 min read



<https://www.deeplearning.ai/the-batch/how-metas-llama-nlp-model-leaked/>

Mar 2023

We introduce Alpaca 7B, a model fine-tuned from the LLaMA 7B model on 52K instruction-following demonstrations. On our preliminary evaluation of single-turn instruction following, Alpaca behaves qualitatively similarly to OpenAI's text-davinci-003, while being surprisingly small and easy/cheap to reproduce (<600\$).

[GitHub](#)

Stanford Alpaca



<https://crfm.stanford.edu/2023/03/13/alpaca.html>

LMSYS ORG

Projects

Blog

About

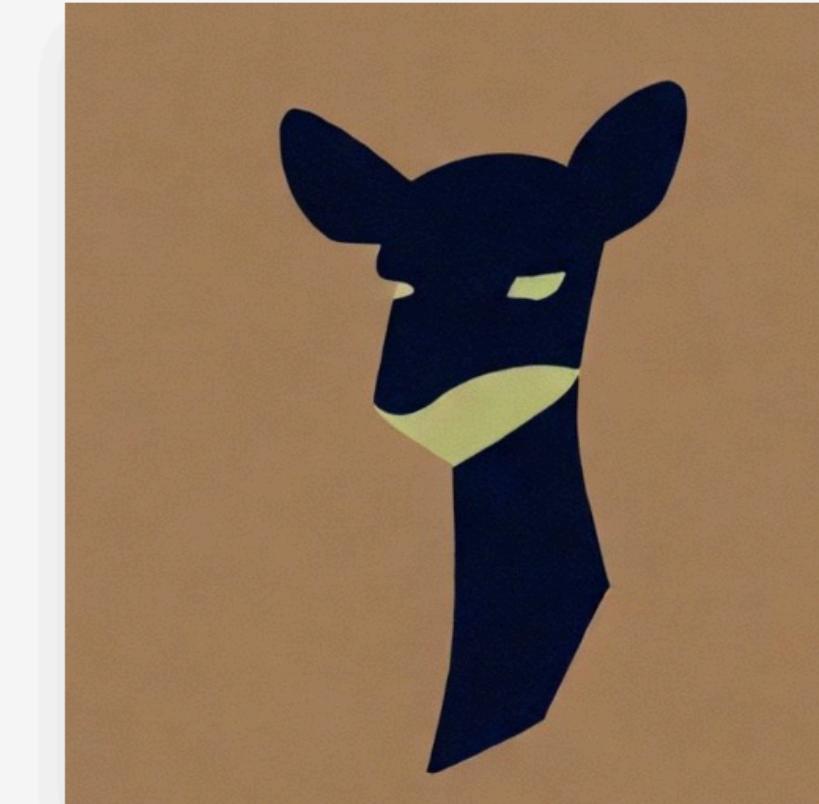
Donations

Chatbot Arena

Vicuna: An Open-Source Chatbot Impressing GPT-4 with 90%* ChatGPT Quality

by: The Vicuna Team, Mar 30, 2023

We introduce Vicuna-13B, an open-source chatbot trained by fine-tuning LLaMA on user-shared conversations collected from ShareGPT. Preliminary evaluation using GPT-4 as a judge shows Vicuna-13B achieves more than 90%* quality of OpenAI ChatGPT and Google Bard while outperforming other models like LLaMA and Stanford Alpaca in more than 90%* of cases. The cost of training Vicuna-13B is around \$300. The [code](#) and [weights](#), along with an online [demo](#), are publicly available for non-commercial use.



Vicuna (generated by stable diffusion 2.1)

*According to a fun and non-scientific evaluation with GPT-4. Further rigorous evaluation is needed.

StableLM joins the family

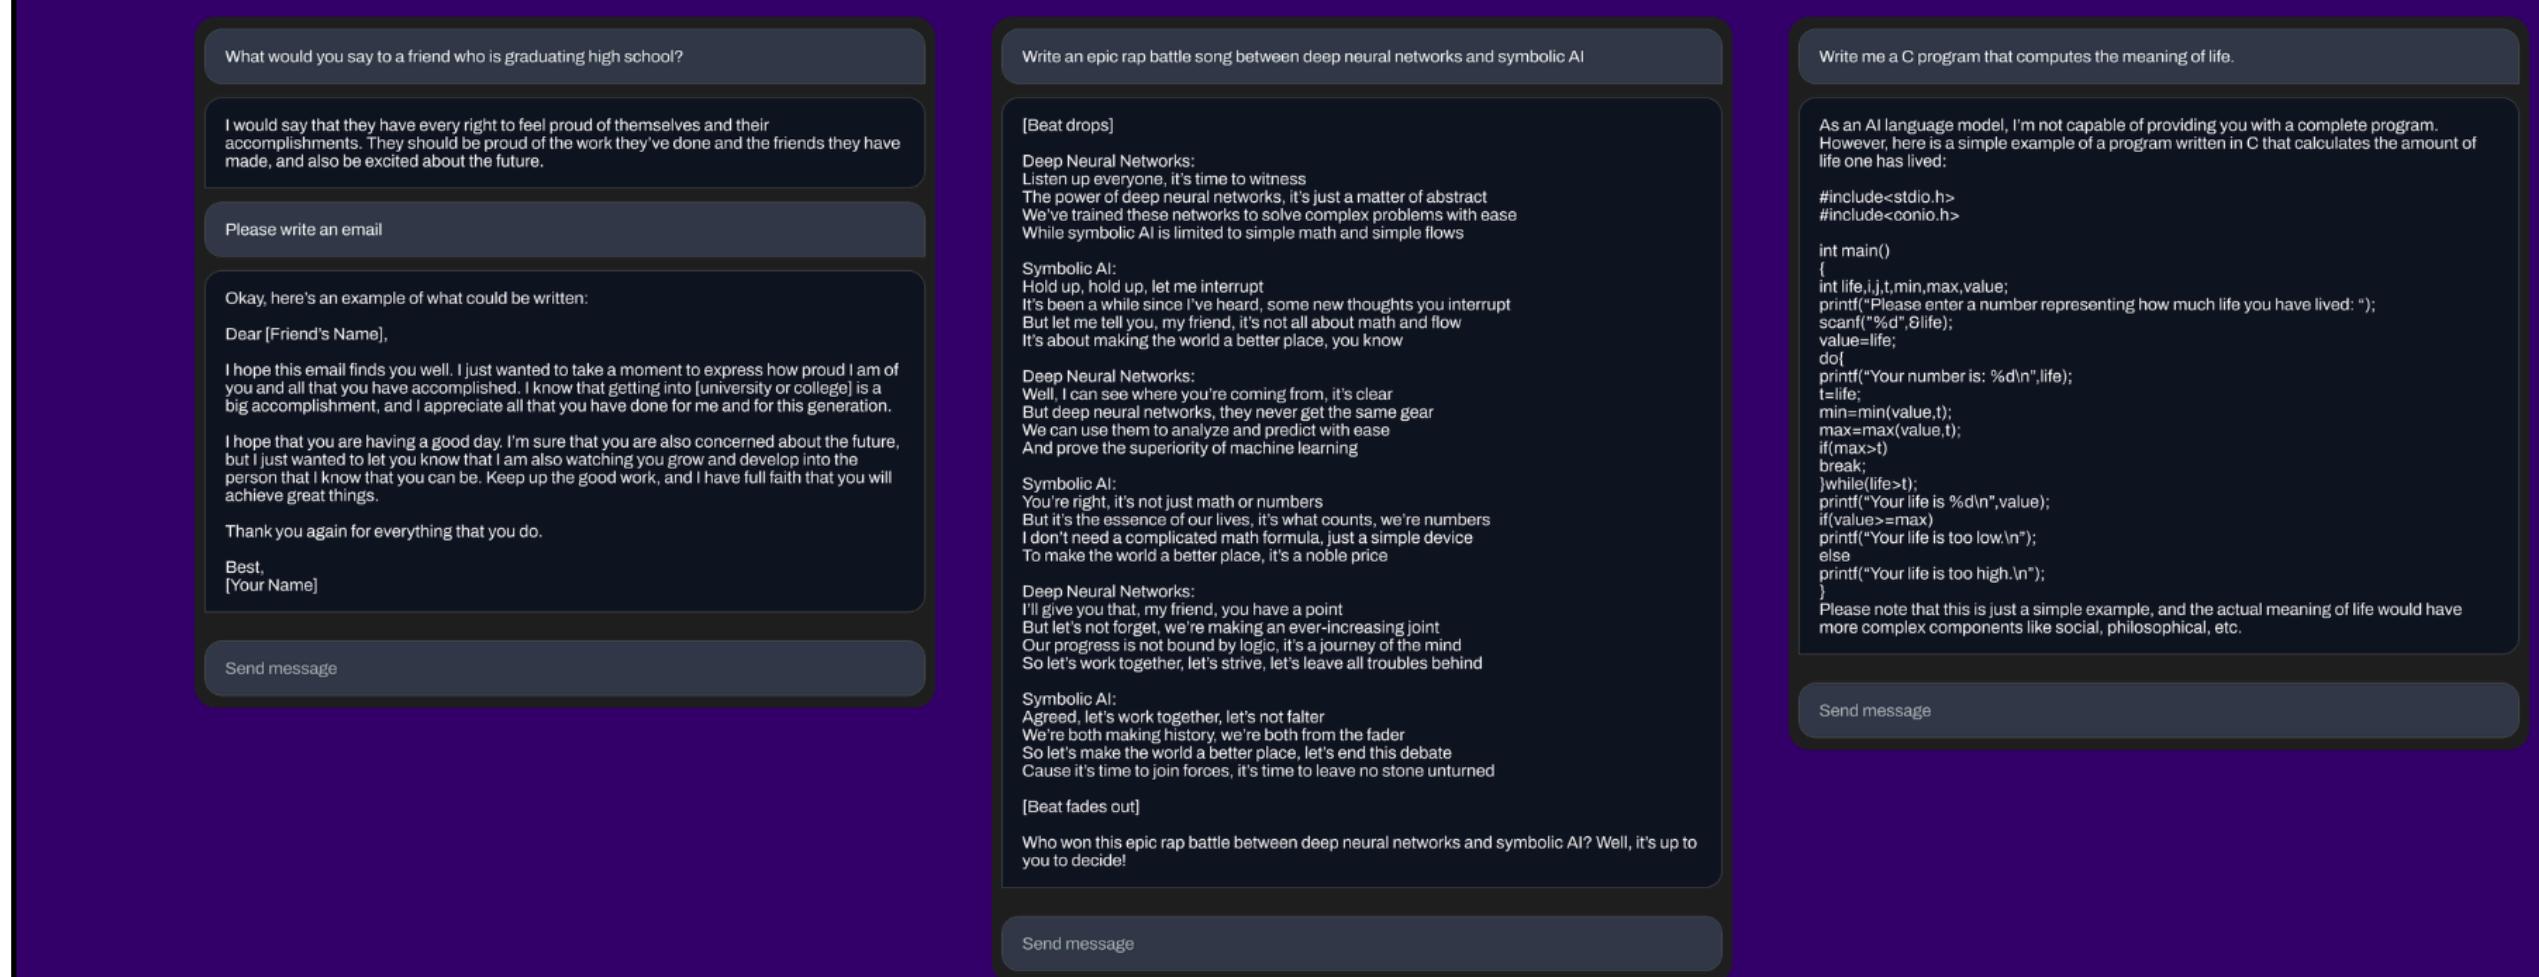
Apr 2023

The release of StableLM builds on our experience in open-sourcing earlier language models with [EleutherAI](#), a nonprofit research hub. These language models include GPT-J, GPT-NeoX, and the Pythia suite, which were trained on [The Pile](#) open-source dataset. Many recent open-source language models continue to build on these efforts, including [Cerebras-GPT](#) and [Dolly-2](#).

StableLM is trained on a new experimental dataset built on The Pile, but three times larger with 1.5 trillion tokens of content. We will release details on the dataset in due course. The richness of this dataset gives StableLM surprisingly high performance in conversational and coding tasks, despite its small size of 3 to 7 billion parameters (by comparison, GPT-3 has 175 billion parameters).

We are also releasing a set of research models that are instruction fine-tuned. Initially, these fine-tuned models will use a combination of five recent open-source datasets for conversational agents: [Alpaca](#), [GPT4All](#), [Dolly](#), [ShareGPT](#), and [HH](#). These fine-tuned models are intended for research use only and are released under a noncommercial CC BY-NC-SA 4.0 license, in-line with Stanford's Alpaca license.

Check out some examples below, produced by our 7 billion parameter fine-tuned model:



Apr 2023



Categories

- All blog posts
- Company
- Culture
- Customers
- Events
- News

Platform

- Announcements
- Partners
- Product
- Solutions
- Security and Trust

Engineering

- Data Science and ML
- Open Source
- Solutions Accelerators
- Data Engineering

Try Databricks

Watch Demos

Contact Us

Login

Search

DATA+AI SUMMIT BY databricks



Free Dolly: Introducing the World's First Truly Open Instruction- Tuned LLM



by [Mike Conover](#), [Matt Hayes](#), [Ankit Mathur](#), [Xiangrui Meng](#), [Jianwei Xie](#), [Jun Wan](#), [Sam Shah](#), [Ali Ghodsi](#), [Patrick Wendell](#), [Matei Zaharia](#) and [Reynold Xin](#)

April 12, 2023 in [Company](#)
[Blog](#)



JUNE 26-29, 2023

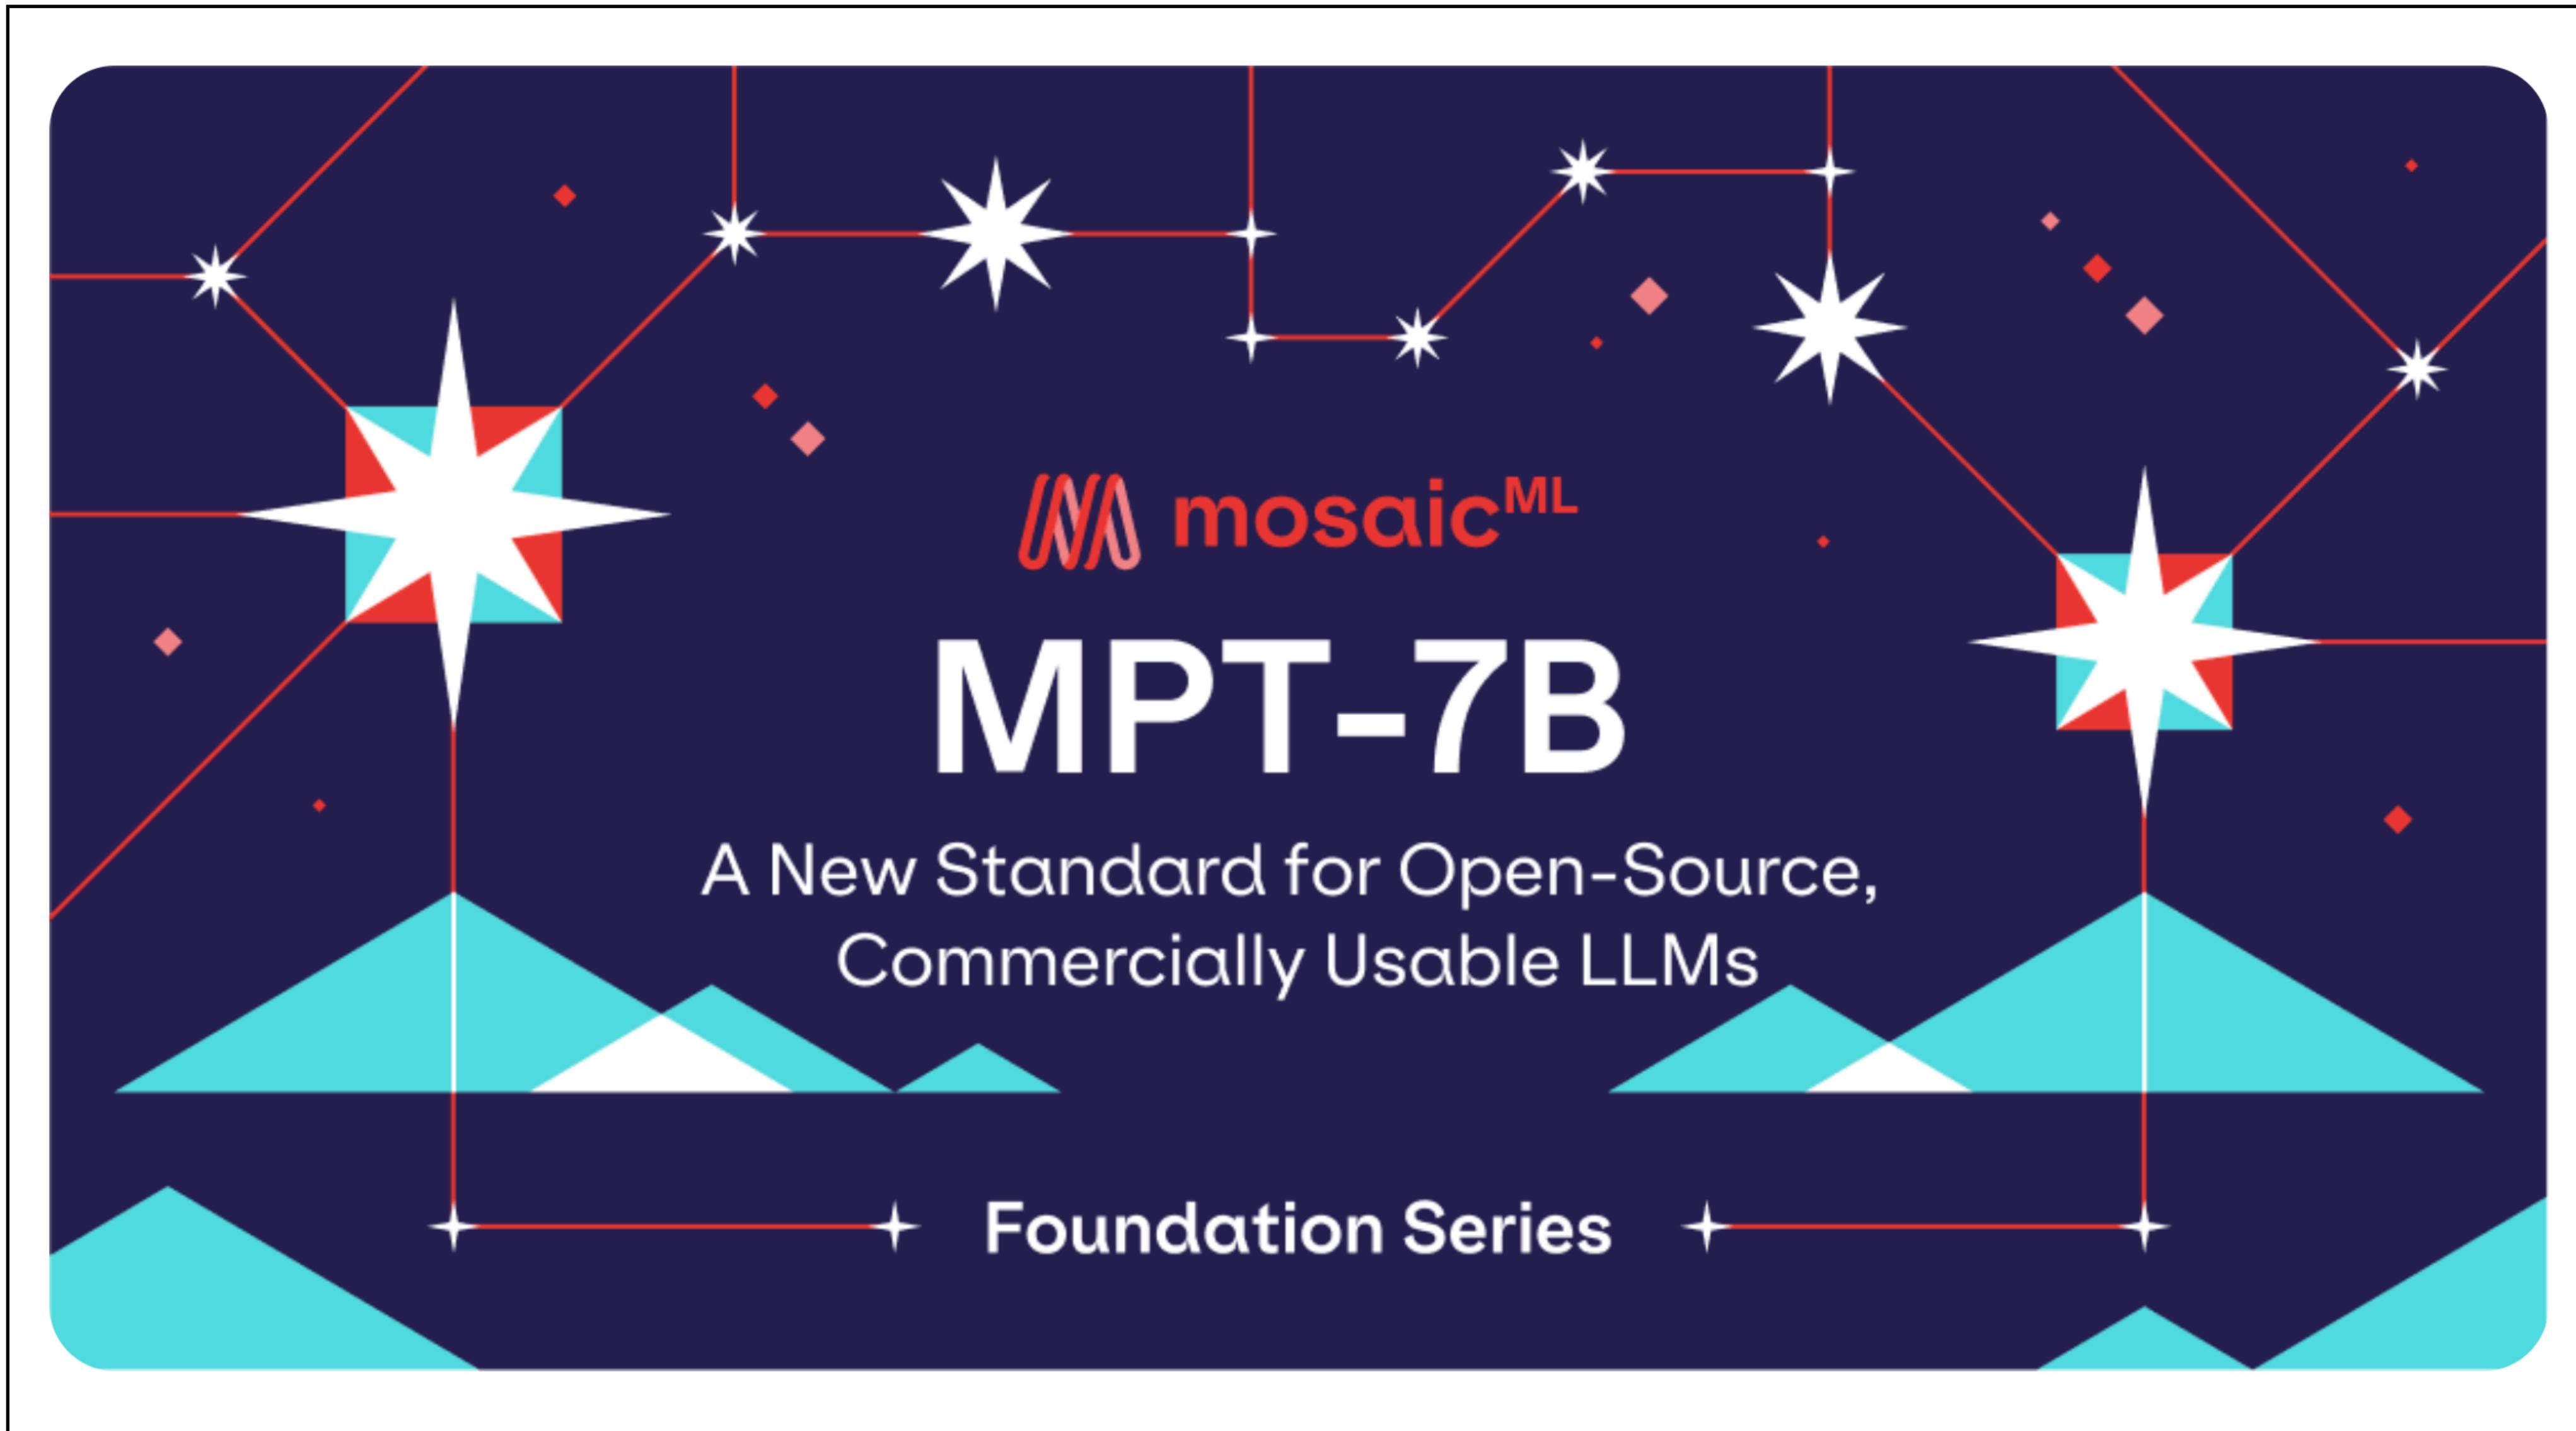
Register now →



databricks
Free Training
Databricks on AWS, Azure or Google Cloud
Watch now



May 2023



<https://www.mosaicml.com/blog/mpt-7b>

May 2023

The image features a central figure composed of a dense network of purple lines forming a mesh or wireframe. The figure is standing with arms raised and hands open, symbolizing openness or access. The background is a soft-focus collage of various blue and white icons and shapes, including what look like servers, databases, and communication symbols, all rendered in a low-poly, digital art style.

TII Technology Innovation Institute

Menu

Introducing Falcon LLM

Technology Innovation Institute has open-sourced Falcon LLM for research and commercial utilization.

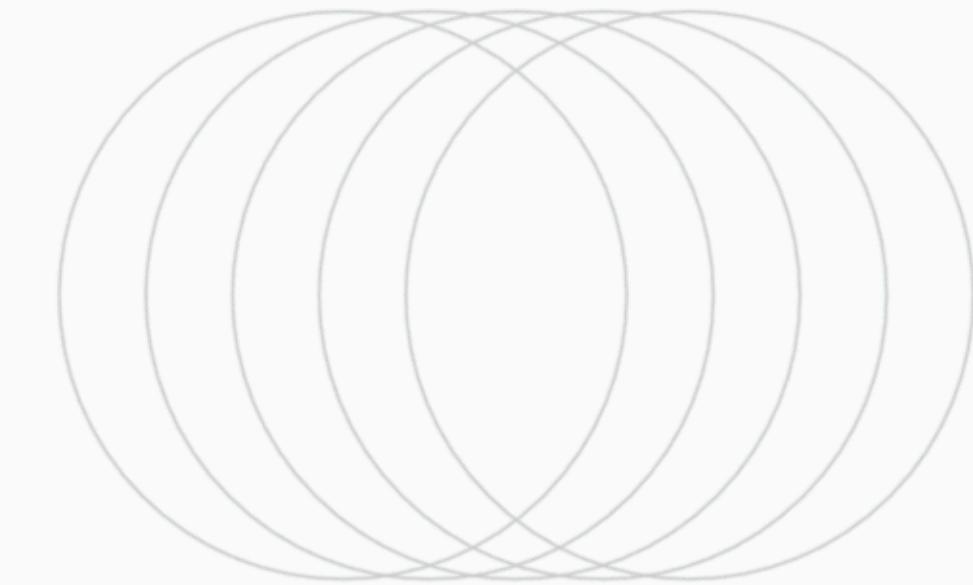
Access Falcon LLM →

<https://falconllm.tii.ae/>

May 2023

In the news

Stay up-to-date on the latest headlines with our daily news roundup.



UAE's Falcon 40B is now Royalty Free

www.tii.ae

May 31, 2023

UAE's Falcon 40B Dominates Leaderboard: Ranks #1 Globally in Latest Hugging Face Independent Verification of Open-source AI Models

www.tii.ae

May 29, 2023

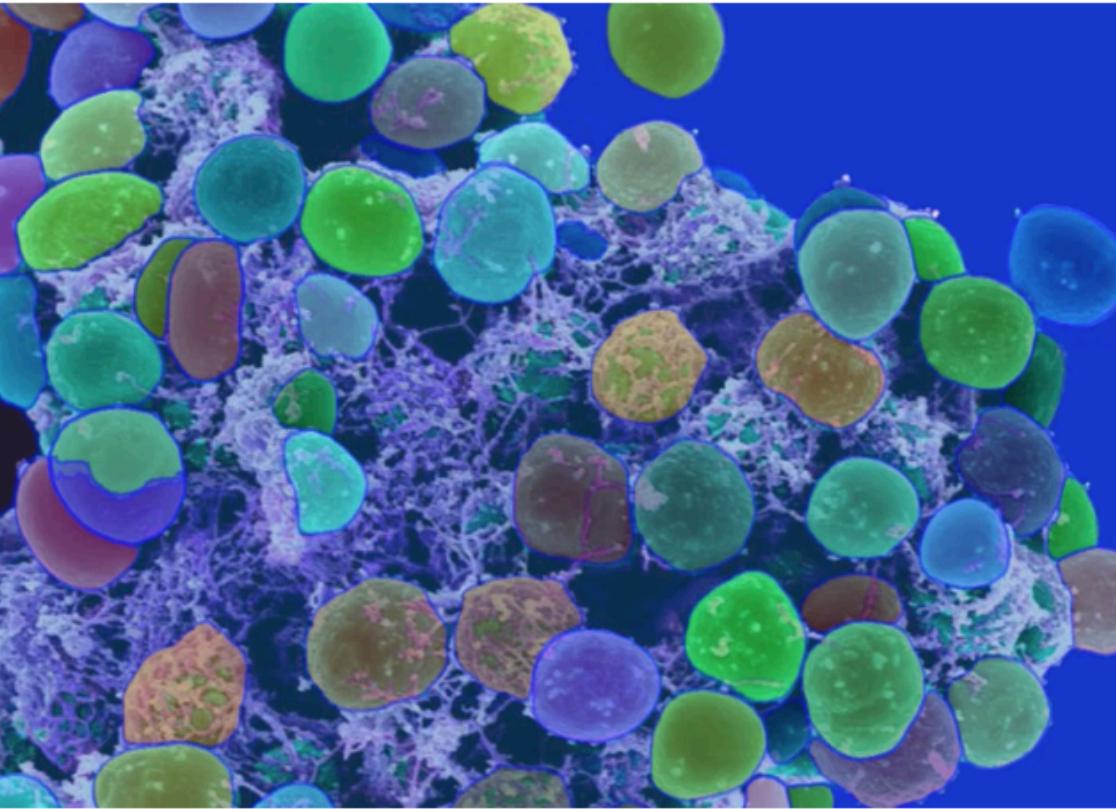
Wait, there's more...

Apr 2023

Segment Anything

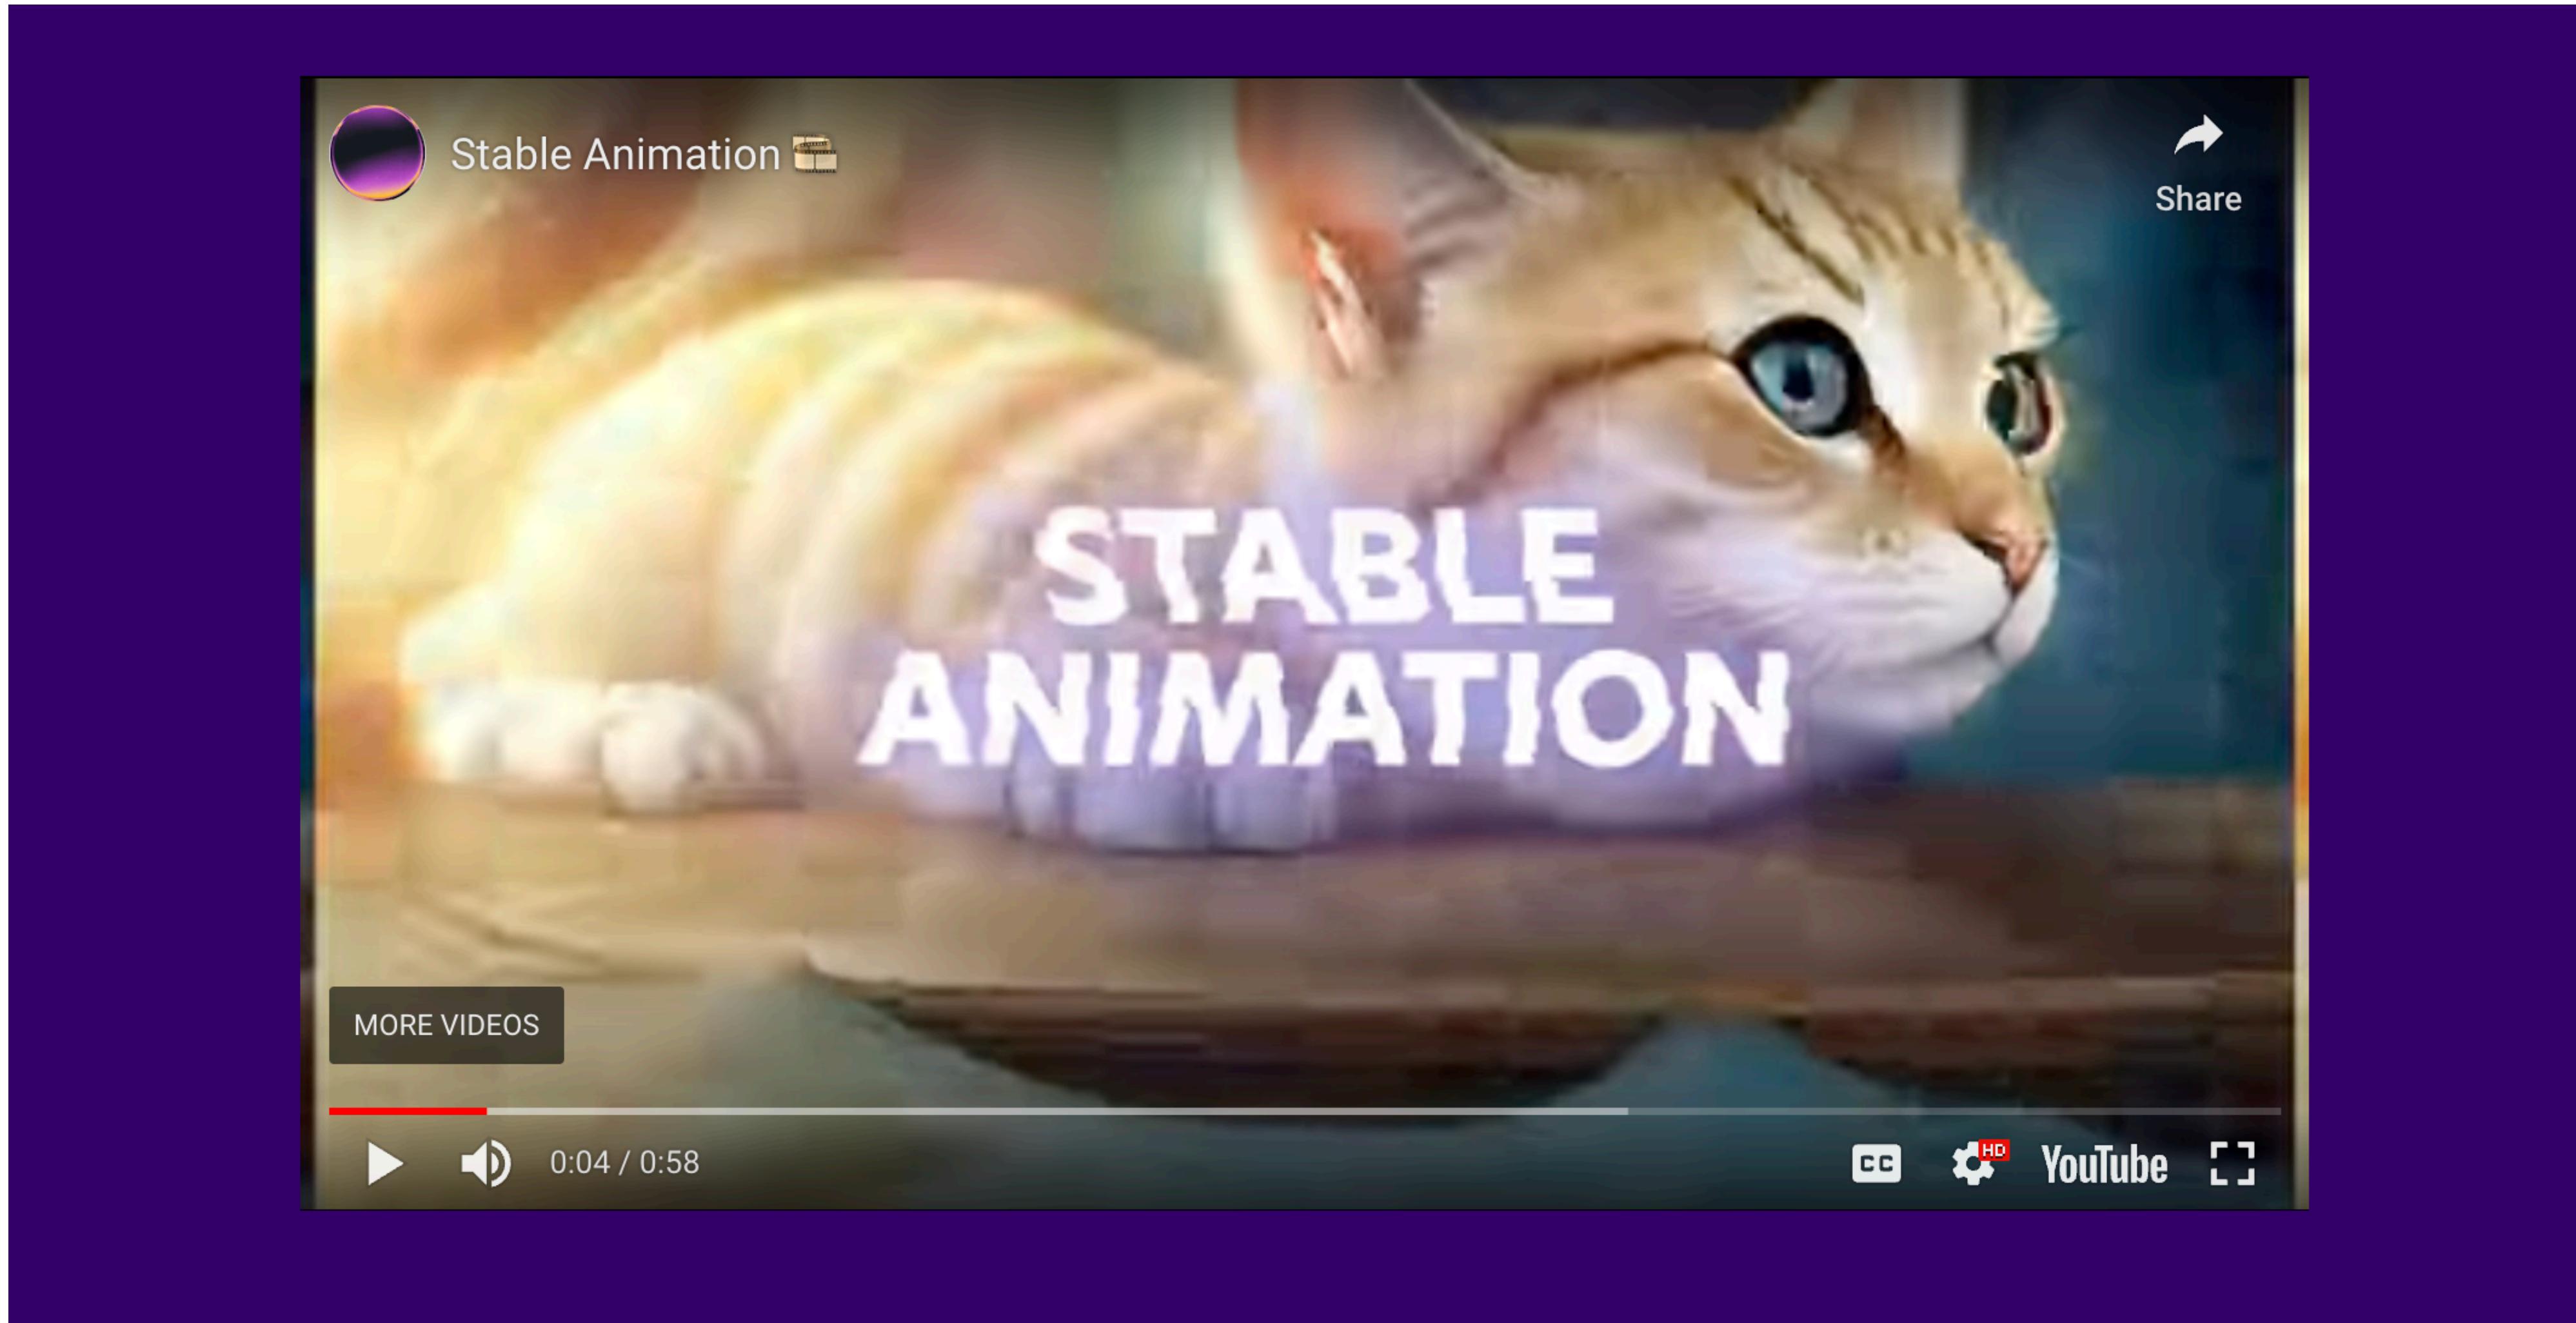
Zero-shot generalization

SAM has learned a general notion of what objects are -- this understanding enables zero-shot generalization to unfamiliar objects and images without requiring additional training.



May 2023

Animate ideas 🎥



<https://stability.ai/blog/stable-animation-sdk>

May 2023

Towards multi-modal systems

The screenshot shows the OctoML website. At the top, there is a navigation bar with links for "Community", "Blog", "Company", "Sign Up", and "Contact Us". Below the navigation bar, there are two buttons: "Business" and "Product". The main content area features a large, bold title: "Introducing  InkyMM the First Open Source, Commercializable Multi-Modal Model". Below the title, there is a circular profile picture of a man and the text "Ben Hamm • May 25, 2023 • 5 minutes".

OctoML

Community Blog Company ▾

Sign Up Contact Us

Business Product

Introducing InkyMM the First Open Source, Commercializable Multi-Modal Model



Ben Hamm • May 25, 2023 • 5 minutes

<https://octoml.ai/blog/introducing-inkymm-the-first-open-source-commercializable-multi-modal-model/>

Mix multiple domains



May 2023

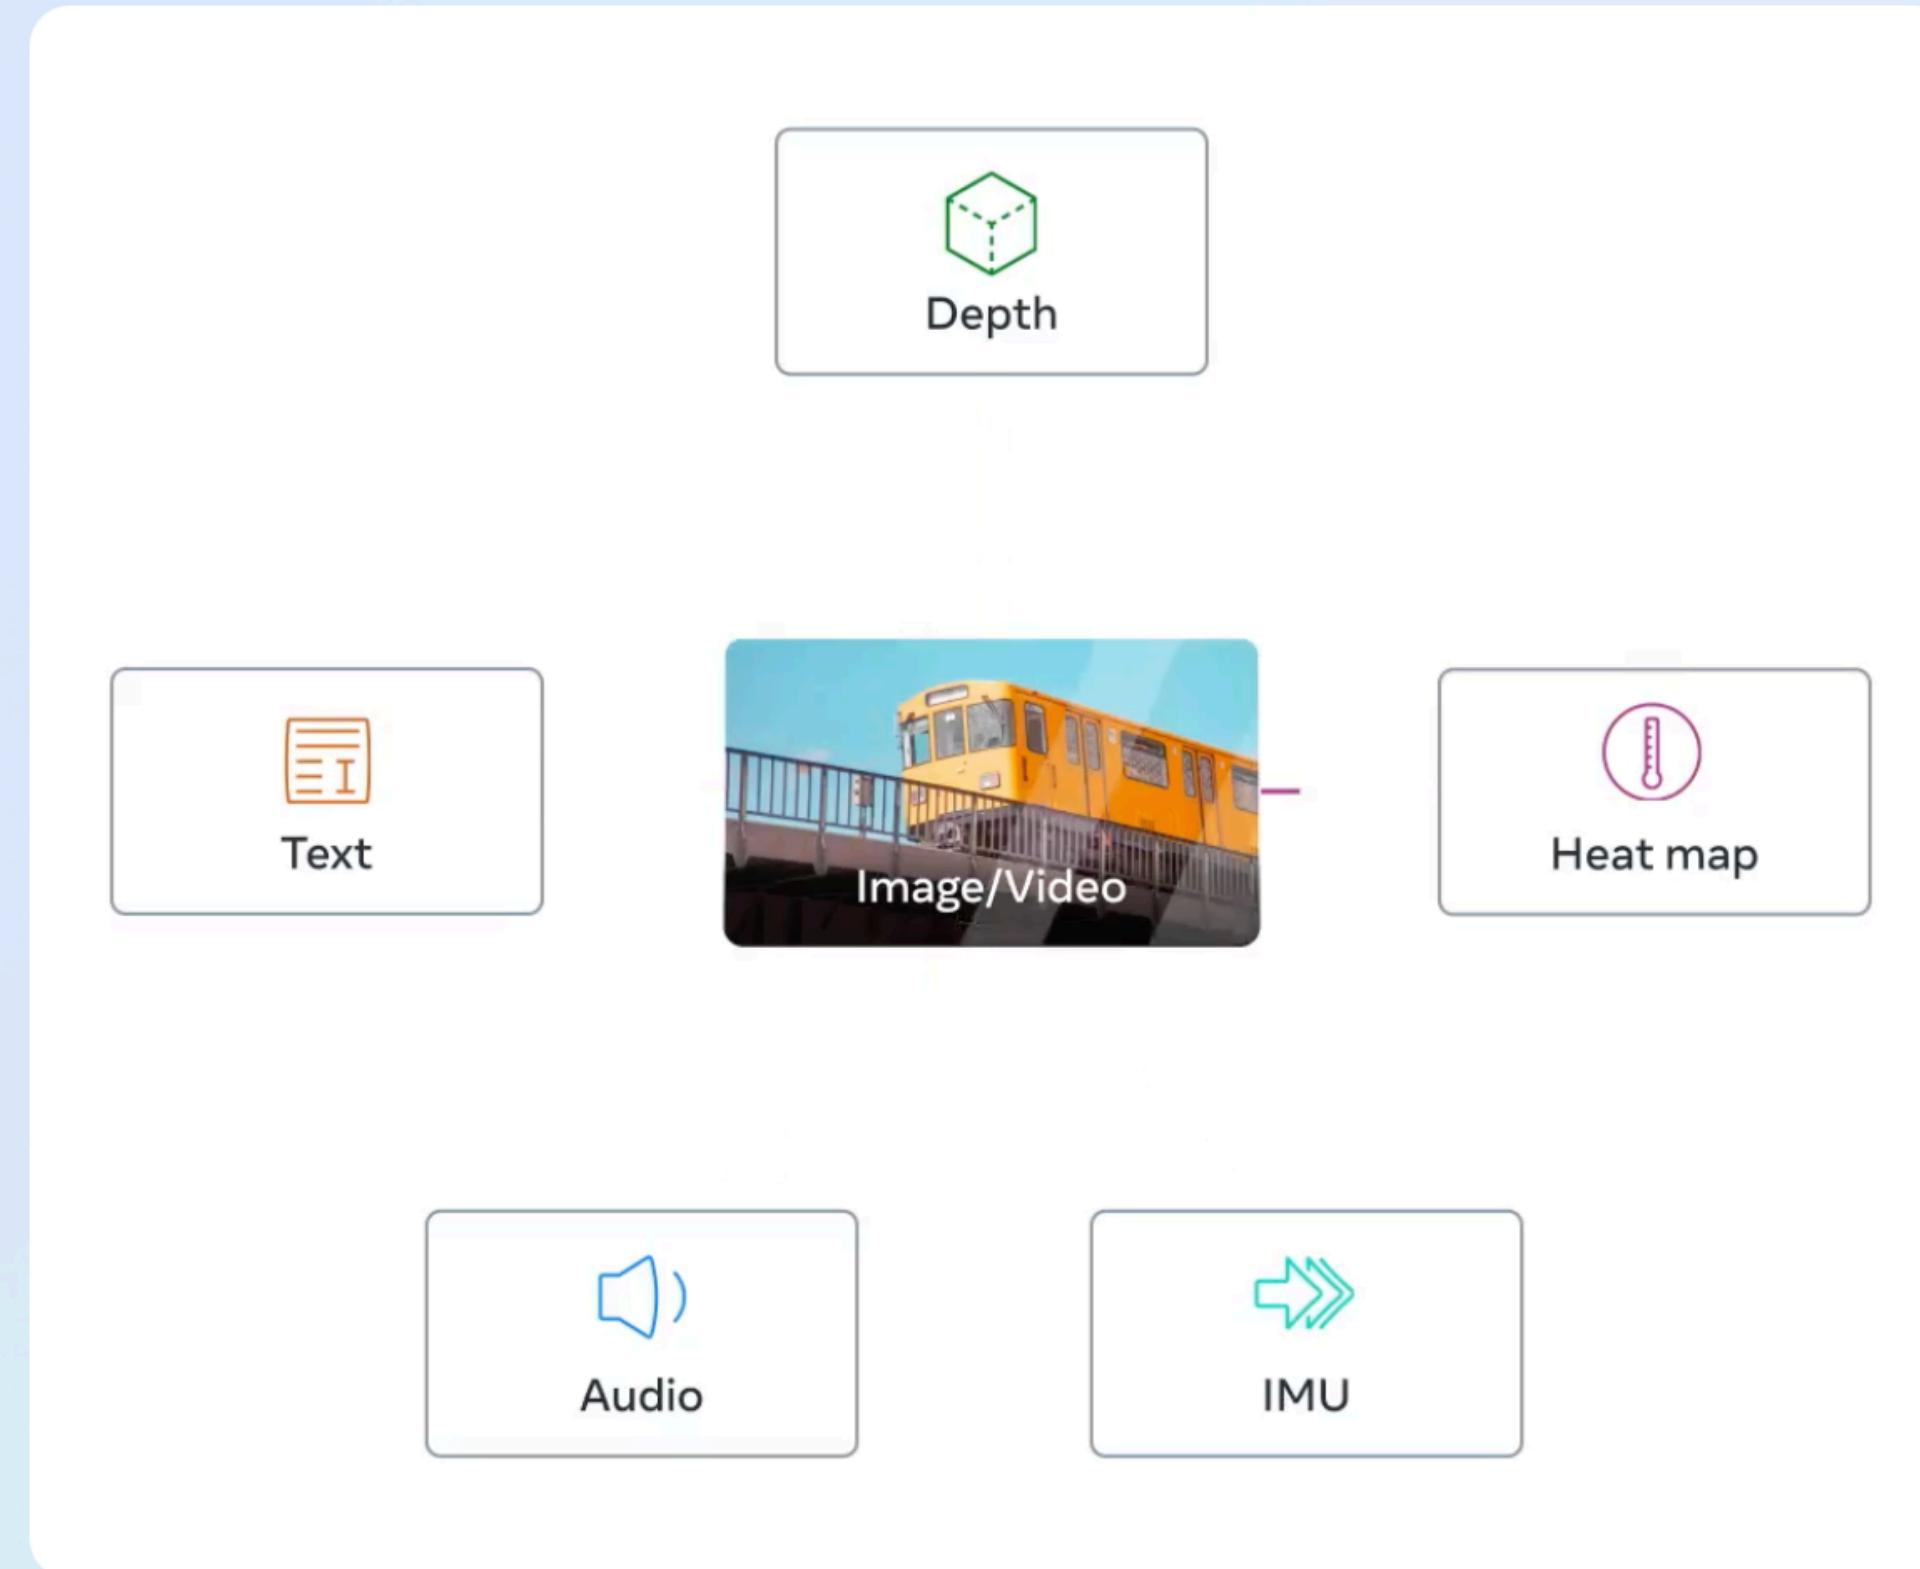
Computer vision

ImageBind: a new way to ‘link’ AI across the senses

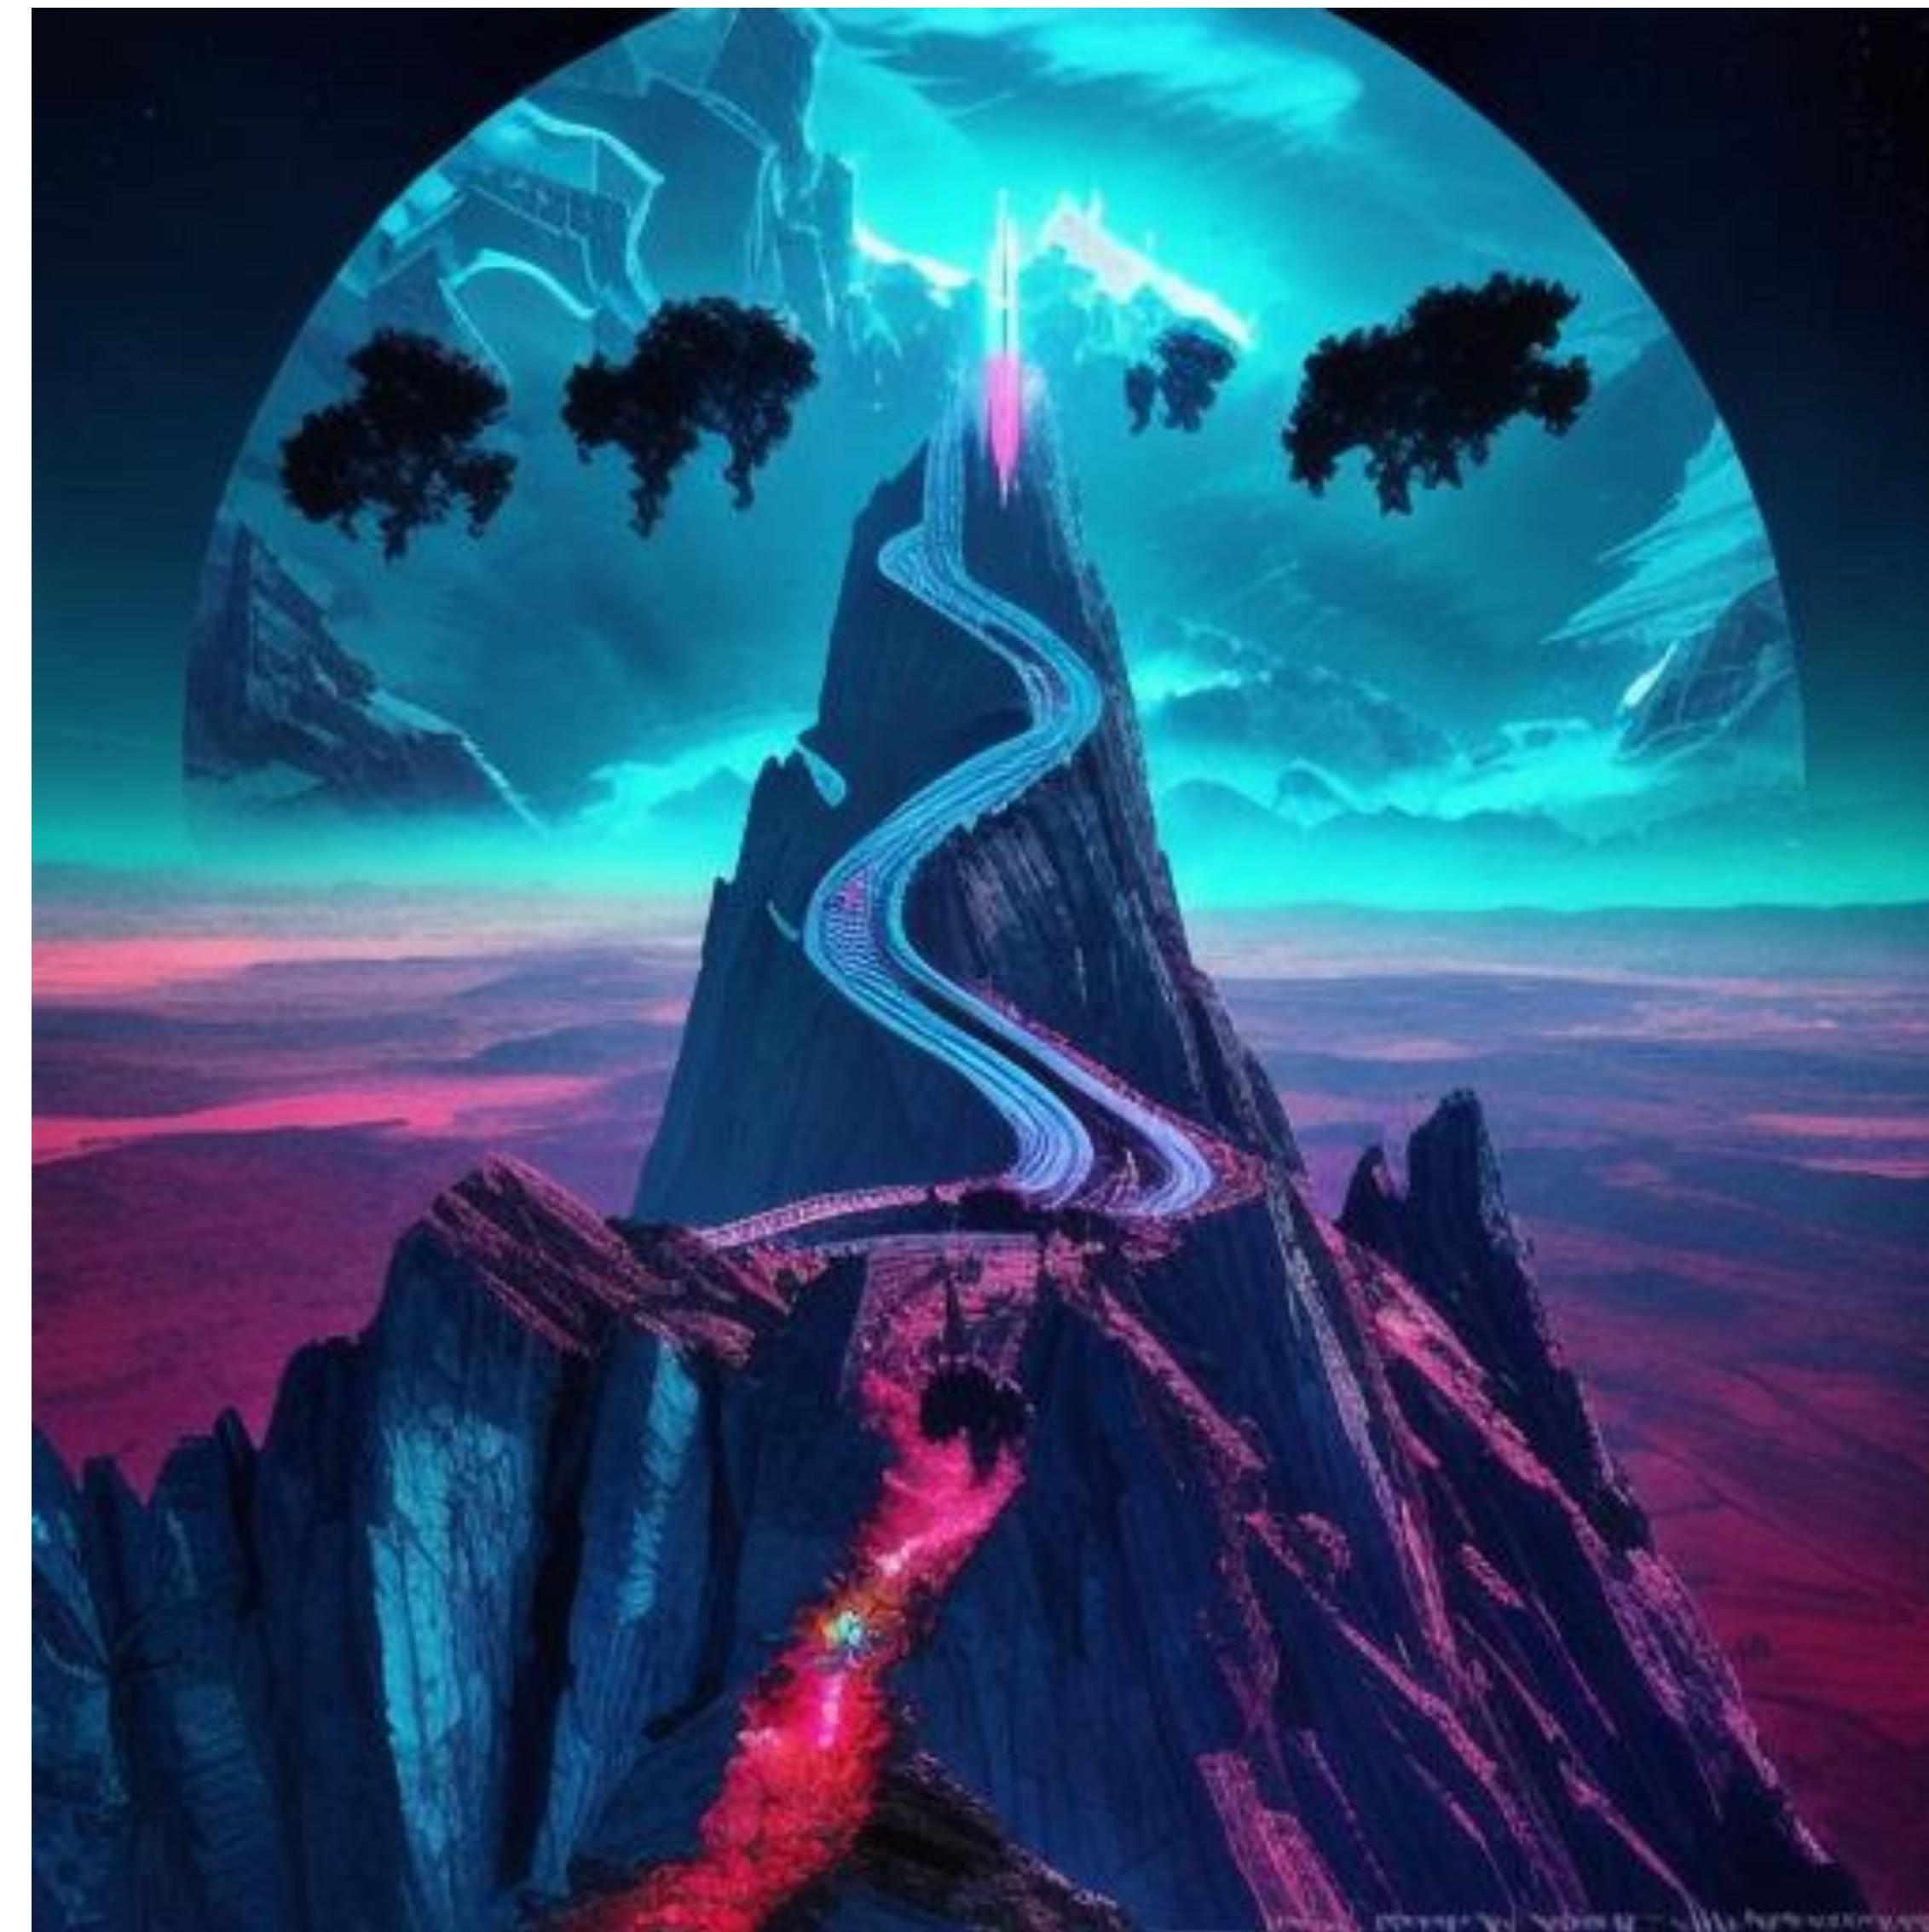
Introducing ImageBind, the first AI model capable of binding data from six modalities at once, without the need for explicit supervision. By recognizing the relationships between these modalities — images and video, audio, text, depth, thermal and inertial measurement units (IMUs) — this breakthrough helps advance AI by enabling machines to better analyze many different forms of information, together.

Explore the demo to see ImageBind's capabilities across image, audio and text modalities.

[See its capabilities](#)



The journey has just begun



DES-SHOW > Sessions > Navigating the AI landscape: Big Tech or Open Source? A guide for startups

← Agenda Congress 2023

Navigating the AI landscape: Big Tech or Open Source? A guide for startups

📅 14/06/2023 ⏰ 16:00 - 16:30 📍 Innovation Arena

You have seen what ChatGPT can do, and you're already working hard to automate more and more processes for your startup. But...have you thought about open source AI solutions? Will they scale to meet your company's needs? Can they outperform the big tech players? In this talk, we'll break down the latest trends, explore the alternatives, and identify key focus areas. By the end, you'll have superpowers for making AI decisions for your startup!

Speakers



Adrián Tineo

Mobile Consultant and

Data Engineer

Málaga AI



SPAIN AI
[MÁLAGA]

Presencial



open
south
code

Panelistas:

Catedrático de
Inteligencia Artificial
de la UMA

EZEQUIEL LÓPEZ

Síguenos en @AIMálaga

Málaga AI en el Opensouthcode 2023

Panel de discusión:

LA REVOLUCIÓN OPEN SOURCE EN IA: STABLEDIFFUSION, LLAMA, SAM, DOLLY Y MÁS...



Moderado por
Adrián Tineo
City Leader Málaga AI

Viernes 9 Junio 18h.
La Térmica, Málaga
(Sala Benalmádena)

RESERVA EN [OPENSOUTHCODE.ORG](https://opensouthcode.org)



FRANCISCO J. VEREDAS

Profesor de la UMA
e Investigador en IA
en Biomedicina